


Figure 6 — A taxonomy of interpretability failure modes on the VCS, by underlying mechanism



Each leaf is a method exhibiting the branch's failure mechanism, annotated with its MEASURED faithfulness F against the §1 intervention oracle (read from leaderboard.json; 0 = chance, 1 = oracle ceiling). Branch (1) is sharpest on the discrete position/index regime, where the naive gradient is provably zero: vanilla saliency F=0.00 there vs F=0.53 on the content regime (a 0.34 causal–gradient bucket gap, faithful_demo.json). The N/A control block is a property of the SUBJECT, not a method failure — Grad-CAM, attention and VIPER need NN substrate the VCS lacks (recorded finding, exp_design §5). Higher-order branches (3–5) name the deepest objections: present≠used probing, non-identifiable factorisations, single-site blind-spots. Mechanism map: plan.md representativeness rebuttal + exp_design §7.