Keystone: the bilinear sampler makes gradient attribution faithful on the position regime — yet recovers no semantics

On the discrete position/index regime the naive gradient is provably zero (Prop. prop:zero), so every gradient method scores faithfulness 0. Turning Paper-1's bilinear index-boundary sampler ON (same surrogate as the SI joystick tool) restores a real position gradient and the methods become FAITHFUL — up to 0.791 (Vanilla saliency on pong). Yet semantic_recovery stays 0 for all 30 (method × game) records in BOTH conditions: an attribution map names no game concept on its own. Faithfulness is repairable; the semantic gap is not — the danger zone is not a fixable artifact of the vanishing gradient.