

Faithful attribution recovers the causal wiring, not the semantics: a ground-truth audit of interpretability on a differentiable microprocessor

Andreas Maier^{1*}, Siming Bayer¹ and Patrick Krauss¹

^{1*}Pattern Recognition Lab, Friedrich-Alexander-Universität
Erlangen-Nürnberg, Erlangen, Germany.

*Corresponding author(s). E-mail(s): andreas.maier@fau.de;

Abstract

Interpretability methods are judged by their faithfulness: whether they name the true causes of a system’s output. The field treats faithfulness as the test that separates a real explanation from a merely plausible one. We test this on the Understanding VCS Project, a bit-exact and differentiable copy of the Atari 2600. The machine was never built to be interpretable. Yet it is fully known, so the true cause of any output can be computed by exact intervention. We score thirty interpretability methods against this known answer. Causal and intervention methods are more faithful than gradient and correlational ones (**0.68** versus **0.29**). The clearest case is a sprite’s position, where the naive gradient is provably zero: a causal method scores **1.000** and standard saliency scores **0.000**. But faithfulness is not what is missing. When we switch on the port’s differentiable sampler, the gradient methods become faithful too—saliency rises from **0.000** to **0.791** on Pong—yet no method recovers what a variable means. Every method reports which signals matter, never what they mean. We name this remaining distance the semantic gap, and we measure it. Meaning is not in the mechanism. For software it is fixed outside the code, in the documentation and the data the program was written for. No method that reads only the mechanism can reach it. The Understanding VCS Project offers a reusable benchmark with known answers, for explainable AI, mechanistic interpretability, and software verification.

Keywords: interpretability, explainable AI, mechanistic interpretability, faithfulness, ground truth, causal attribution, semantic gap, Atari VCS

1 Introduction

Interpretability now matters in practice. Neural networks drive cars, read scans, and approve loans. Before we trust such a decision, we ask why. Yet on the systems we most want to explain we cannot check whether the answer is correct. A saliency map for a deep network can look *convincing* and can change with the action. That does not mean it delivers a *correct explanation* [1]. So far, the field of explainable AI has converged on *faithfulness* as its quality criterion, that is, whether an explanation names a system’s *true* causes. The difficulty is that on a real neural network there is nothing to measure it against.

Similar shortcomings were also identified in neuroscience. Jonas and Kording showed this clearly. They applied standard neuroscience methods to a microprocessor in which every transistor is known [2]. Lesions, tuning curves, and connectomics produced rich structure, but none of it recovered how the chip computes. Without ground truth we cannot distinguish a method that finds some structure from one that finds the *correct* structure. The same blind spot recurs across explainable AI, where attribution maps are validated against each other or against human intuition rather than against the system’s mechanism [3]. Faithfulness has been used as a substitute for understanding, but no one could test this substitute.

Maier et al. recently built the missing reference. It is a bit-exact, differentiable port of the Atari 2600 Video Computer System (VCS) that reproduces a real 1977 machine to the byte and pixel [4, 5]. The VCS is software with a known specification, so the true cause of any output is computable by exact intervention. This *intervention oracle* (cf. Section 3.2) returns the true causal effect for any candidate cause of any output, measured rather than estimated. It operates over the platform shown in Fig. 1. On this machine an interpretability claim can be scored right or wrong.

We study three traditions: neuroscience, attribution, and mechanistic interpretability. This covers 30 interpretability methods plus the oracle positive control, across 31 experiments and 257 per-game records. We score each method on two axes: faithfulness against the oracle, and a documented plausibility proxy for how convincing it looks. The first result is the expected one. Causal and intervention methods are faithful, while gradient and correlational saliency methods are not. The contrast is sharpest on discrete sprite-position outputs. There the naive gradient is provably zero, because position is set by strobe timing rather than by a value it can flow through. In that position regime activation patching reaches faithfulness 1.000 while vanilla saliency collapses to 0.000 on all six core games. The saliency method nonetheless has the higher plausibility score (0.9 versus 0.5). It looks most convincing exactly where it is least correct. This is the region Fig. 4 marks as conviction without correctness. This is the result the field would predict, and we treat it as an intermediate step, not the main finding.

Faithfulness is not what separates good explanations from bad ones. A reader could dismiss the zero of the gradient methods as a fixable artifact of a discrete index. So we test the strongest version of the question and close this objection. We turn the differentiable port’s bilinear sampler on. The sampler interpolates the position read and restores a non-zero gradient through it, and the gradient methods become faithful too. On the games with a moving tracked sprite their faithfulness rises off

the provable zero. Vanilla saliency, $\text{grad} \times \text{input}$, smoothgrad, and expected gradients reach 0.791 on Pong, and the same family reaches 0.529 on Q*bert (committed record `sampler_faithfulness.csv`, Fig. 5). The increase is real but does not occur on every game. Integrated gradients with a zeros baseline still vanishes on Pong. A static-sprite game offers no position to restore, and where the tracked cell is not the active one nothing moves. Yet on every game where the gradient becomes faithful the semantic recovery stays exactly 0, the same 0 every other method scores, faithful or not. The restored gradient concentrates on the position RAM byte that drives the sprite, `ram[54]` on Pong. It identifies the correct cause but gives no meaning for it.

Securing faithfulness is necessary, but it does not solve the main problem. Faithfulness is necessary for understanding but far from sufficient. Making it universal does not close the gap, because the gap is not a faithfulness deficit. With the same oracle we ask not only *which* methods are faithful but *what* a faithful method delivers. The answer is always the same kind of object: a true causal-effect map or data-flow graph. It is never the variable’s meaning and never the routine’s algorithm. No method on this fully observable, fully differentiable machine names “missile,” “collision,” or “game restart.” Every method returns the same wiring fact: change input X and outputs X and Y move. To make this precise we call an explanation right only if it is **Faithful** (its claimed causes are the true ones), **Sufficient** (it regenerates behavior under held-out interventions), and **Minimal** at the algorithmic level (right = $\mathbf{F} \wedge \mathbf{S} \wedge \mathbf{M}$, Section 3.4). The committed data show the gap in both directions. An automatically discovered circuit on Breakout is graph-correct and minimal yet does not reproduce the behavior it explains ($\mathbf{F} = 1.0$, $\mathbf{M} = 1.0$, $\mathbf{S} = 0.44$ on Breakout). A correct wiring diagram is not the computed function. A learned feature dictionary on Pong matches a named game variable perfectly (matched fraction 1.0) while being causally almost inert ($\mathbf{F} = 0.04$ on Pong). Looking like a known variable is not the same as computing it the way the program does. A concept can be linearly decodable yet provably unused. Being present in the state does not mean being used. On this machine we can measure the gap between a faithful attribution and an account of the computation as a number, which we name the *semantic gap*. It does not shrink when faithfulness is made universal.

The reference for meaning, then, lies outside the mechanism. Every semantic label in the study entered from outside the system. It was imported from external annotation sets [6, 7] and only *causally verified* on our substrate, never *discovered* by a method under test. An alignment method that reports perfect accuracy is therefore right by construction, since it checks a variable we supplied. Probing shows that information is present, not that it is used [8], and our oracle is the strictly stronger test. What fixes the meaning is the system’s *behavior*. For software, that behavior is set down in its documentation and data: the manuals, the specification, and the design. IEEE 610.12 makes these part of the definition of software itself [9]. The wiring is in the mechanism, so any faithful method can recover it. The meaning is not, and no method we tested supplied it.

A 1977 chip may seem the wrong place to judge methods built for learned networks. We read the VCS as a *necessary-condition screen*, not a predictor of success. Here we have perfect ground truth and no learning obscuring the mechanism. Failure under these conditions is strong negative evidence for the claim that the same method reliably

recovers causal mechanisms in harder, less observable systems. The VCS still exhibits the properties that make interpretation hard, from aliased RAM cells to race-the-beam timing. What it does not capture we leave to later work on learned agents.

Because the central claim is sharp, we are deliberate about what we do not claim. We do *not* claim that any method recovered the program’s meaning, algorithm, or design from scratch, and the plausibility axis is a documented proxy, not a measured human study. As such, this paper *defines and measures* the semantic gap and locates its reference in the system’s behavior. *Closing* the gap means recovering a system’s documentation and design from the system alone. That is companion work, not this paper.

The study thus offers a reusable ground-truth benchmark for explainable AI, mechanistic interpretability, AI safety, and software verification, and it reframes the open problem. The hard part of interpretability is not faithfulness. It is recovering meaning, which remains unsolved even when faithfulness is achieved. We believe that naming and measuring this gap is a first step toward closing it.

2 Related work

Interpretability is studied in three communities that rarely cite one another. Yet all ask the same question: inside a complex system, what caused this output? The neuroscience tradition records and perturbs a system, through tuning curves, lesions, connectomics, and directed influence. Jonas and Kording [2] applied these methods to a microprocessor running classic games. They found that it produces rich, publishable structure while missing how the chip computes. Their study reported what the analyses found, but not how much each analysis deviated from the true mechanism. We supply that number (Sec. 4).

The attribution tradition asks, for a single decision, which inputs mattered. Gradient and saliency maps [10–13], class-activation methods [14, 15], occlusion [16–18], and surrogate or game-theoretic attributions [19, 20] form the toolkit, applied to Atari agents to visualise what a policy attends [1, 21]. Yet a counterfactual analysis of those same maps found them exploratory, not explanatory [3]: the highlighted regions need not be the causes. On a learned policy the true cause is unknown, so this doubt can only be argued. We resolve it on a system where the true cause is computable (Sec. 5).

Mechanistic interpretability, the youngest tradition, aims past attribution at the circuit, the subgraph that implements a behaviour. Path patching, causal mediation, and scrubbing intervene on internal activations [22–25]; alignment search and interchange interventions test whether a hypothesised variable is causally realised [26, 27]; sparse dictionaries and probes read features off the state [8, 28–31]. On a real artifact these prove the most faithful (Sec. 6). Yet the same oracle shows that a recovered circuit is a correct data-flow graph. It does not give the function the circuit computes. The sparse-dictionary branch is itself under scrutiny. SAEBench [32] reports that gains on the proxy metrics used to tune sparse autoencoders do not reliably translate into practical interpretability. We probe the same concern from a different side, asking on a known system whether a recovered feature is causally used rather than only decodable (Sec. 6).

The agreed criterion across these traditions is faithfulness: claimed causes should be the true causes. Jacovi and Goldberg [33] make it precise, yet its diagnostics stay indirect when ground truth is absent. Adebayo et al. [34] sharpened this into the sanity check: a saliency method should change when weights or labels are randomised, and several popular methods do not. But such checks are only negative. With no reference map, they cannot certify that a method is correct. So the field optimises a proxy it cannot directly score. Our oracle supplies that map and converts faithfulness into a measured quantity: on the discrete sprite-position regime `activation_patching` reaches faithfulness 1.000 while the popular `vanilla_saliency` collapses to 0.000 on all six core games [35], and the leaderboard plots all 31 methods against the same ground truth (Fig. 4) [36].

The concepts we score against are external. AtariARI labels object positions, scores, and lives from commented disassemblies as a probing benchmark [6]; OCArari adds offset corrections so the coordinates align with the render [7]. These labels are imported, not discovered. We verify them causally by intervention: set the byte, re-render, confirm the object moves to its predicted place on the bit-exact framebuffer. This step states our honesty contract. Every concept-level (T3) score measures alignment to a supplied label. It does not measure a meaning the method found on its own. Note that a concept can be linearly decodable yet play no causal role. Control tasks were built to catch this “present $\neq$ used” trap [8]. Our intervention test is stronger than probing, and that separation is one of the three counted in Fig. 6.

Ours is not the first study to score interpretability against a known answer, yet existing benchmarks build the answer into the system. Yang and Kim’s BAM [37] grades attribution maps against images with known relative feature importance, and M4 [38] extends this idea into a faithfulness benchmark spanning metrics, modalities, and models; in both, the feature ground truth is synthetic, installed by the data-generation or training procedure. Tracr [39] compiles programs into transformers with circuits known by construction; InterpBench [40] trains semi-synthetic transformers to a specified circuit. The Mechanistic Interpretability Benchmark [41] pushes furthest in our direction, scoring circuit localization and causal-variable localization against curated answers in neural language models. Yet each of these answers is engineered for the benchmark, whereas our subject is a real, hand-authored artifact whose ground truth is independent of it. The VCS was never designed to be interpretable. Even so, an exhaustive intervention oracle can recover its full causal structure, on a bit-exact, differentiable port [4], built on the Atari substrate [42] and the reference emulator [5]. Our novelty is not the Atari environment, which the Arcade Learning Environment and its descendants [42, 43] already supply, and which faster and object-centric variants such as CuLE [44], OCArari and JAXAtari [7, 45] extend. The novelty is the bit-exact, differentiable intervention oracle that turns a known machine into a ground truth for interpretability. More importantly, the engineered benchmarks validate faithfulness only, because in them the circuit is defined to be the meaning. On a real artifact the meaning is a separate object from the causal structure. This lets us separate faithfulness from understanding. We can show that a method which perfectly recovers the true wiring still has not recovered the semantics (Fig. 6). We measure that remaining gap, which we call the semantic gap. Closing it means recovering a program’s variables

and design from scratch. We are explicit that this is a different problem, the subject of companion work. To our knowledge, this is the first cross-tradition audit of these methods on a real hand-authored software artifact with an exact intervention oracle; we claim no unqualified “first ground-truth calibration,” since synthetic known-answer benchmarks and Jonas and Kording’s microprocessor study [2] precede us.

Our use of “interpretable” is one of several in circulation. We operationalise it as ground-truth recovery: an explanation is interpretable to the degree that its claimed structure matches the system’s true causal structure, scored by the oracle. Barbiero et al. [46] instead define it as inference equivariance, a design target for new models, while ours audits a system already known. Marr’s levels of analysis [47] and Lazebnik’s “fix-the-radio” parable [48] support this paper’s main point. A faithful description of a system’s wiring is not yet an account of the algorithm it runs. On a known machine, we measure that gap.

3 Methods

Our method rests on a single idea. We study an interpretability question on a machine whose answer we can compute exactly. We describe in turn the substrate that makes the ground truth obtainable, the oracle that delivers it, the tiered semantics we attach to it, and the correctness triad that scores understanding as numbers. The per-method setups live in the Supplementary Information. Fig. 1 sketches the pipeline.

3.1 The differentiable, bit-exact VCS substrate

The subject of this study is the Atari Video Computer System (VCS) itself, not a learned agent. It is a real, complex artifact never designed to be interpretable. Yet we can measure its every causal dependency exactly. Paper 1 built it, a bit-exact, differentiable port of the 1977 VCS. The port is validated bit-for-bit against the `xitari` C++ reference on all 64 games of the Paper-1 roster, in random-access memory and rendered screen, within the validated conformance window [4, 5]. We use it in two interchangeable, cross-checked realizations, `jutari` (Julia) and `jaxtari` (JAX/Python).

The substrate grants four capabilities that no neural network offers for free. The first is an exact forward map,

$$y = \text{VCS}(\text{ROM}, \text{inputs}[0:t], s_0), \quad (1)$$

that is, the output y at step t , a pixel, a score, or a game event, is a deterministic function of the cartridge bytes, the inputs up to t , and the initial state s_0 . Re-running reproduces y to the bit. The second is exact interventions. We set any cause u , a ROM byte, a RAM cell, a CPU/TIA/RIOT register, an opcode, or a joystick input, to any value and re-run Eq. (1) deterministically. This realizes the causal operator $\text{do}(u)$ on a real machine rather than an assumed world model. The third is full state observability. We record every bit of RAM and every register along the trajectory, so the “activations” we interpret are the program’s own variables. The fourth is a content-path gradient $\partial y / \partial u$ through a straight-through estimator (STE) formulation

A differentiable, bit-exact game console as a ground-truth testbed for interpretability

The subject IS the substrate — a fully known program — so every explanation can be scored against the true causes.




Fig. 1 The measurement platform and its ground-truth oracle. A differentiable, bit-exact port of the Atari VCS (left), realized as `jutari` (Julia) · `jaxtari` (JAX), gives four capabilities no neural network grants for free: exact forward, exact interventions `do(u)`, full state observability, and content-path gradients. On it we build the oracle (centre): the primary intervention oracle of Section 3.2 (Section 3.2), $\text{do}(u:=v') \rightarrow \Delta y(u)$ by bit-exact re-run, with a content-path gradient $\partial y / \partial u$ as a companion on the content path only. Any explanation $\hat{E}$ is scored by the correctness triad (right), $\text{right} = F \wedge S \wedge M$, and reported on two axes, faithfulness against the oracle versus a plausibility proxy. Faithfulness aggregates correlation, sign and magnitude agreement, deletion/insertion area-under-curve, and precision at k against the true causal top- k (Section 3.4). A discrete index — a sprite’s position, set by strobe timing — has a naive gradient that is zero under the hard forward map (Prop. 1); for position, index, and event outputs the intervention oracle is the *sole* ground truth there: gradient methods are scored as methods under test. The schematic draws no experiment data. Data and oracle: see Supplement; oracle defined in Section 3.2 (Section 3.2).

whose forward pass is bit-exact to the hard machine at any finite temperature [49, 50], batched on the GPU. A neural interpretability study works with three things: the inputs, the state, and the outputs. Our substrate gives us all three.

We respect one limit throughout. Paper-1 bit-exactness holds within a bounded frame window on a fixed action stream, on the order of 30 frames of RAM and 60 frames of screen. Every scored claim is taken inside that validated horizon. Long-horizon end-to-end gradients are not needed for our single-step metrics, so we do not assume them.

3.2 The ground-truth oracle

The oracle is central to the method. For any output y and cause u , it returns the true effect of u on y . We score every method against it. Its primary construction is the exact intervention oracle: we intervene on u by occlusion, clamping, replacement, or

resampling, re-run the machine deterministically, and record the change in the output,

$$\Delta y(u) = \text{VCS}(\dots, \text{do}(u:=v'), \dots) - y. \quad (2)$$

The true causal effect of a cause is the difference its edit makes to the output. We measure it by running the altered machine. We assume no world model and no smoothness. The collection $\{\Delta y(u)\}$ over all causes is the true causal map of y , the reference for the faithfulness axis of every method.

The output y comes in five regimes. The regime fixes what a method must predict and whether a gradient is even meaningful. A *content* output is a register, color, or graphics-bit value that flows smoothly into the pixel, where the naive gradient is valid. A *position/index* output is a sprite, missile, or ball coordinate set by strobe timing. Here the naive gradient vanishes by Prop. 1, and only the bilinear-sampler surrogate or the intervention oracle applies. An *event* output is a discrete state change, a collision latch or a life lost, also gradient-zero. A *score* output is the game’s running counter, valid for the gradient only where it varies continuously. A *pixel* output is the rendered framebuffer value. It is content-valid where it tracks a smooth cause and index-zero where it tracks a position. The Supplement lists, per game, which outputs each method is scored on and the per-output gradient verdict.

The oracle uses several kinds of intervention. We list each one, because the kind limits what a score means. Occlusion and clamping are surgical but *off-manifold*: they can set a byte to a value the running program would never write. Replacement from a donor trajectory is *on-manifold*: the substituted value is one the machine actually produced elsewhere. Resampling draws the cause from its own marginal over the trajectory. Counterfactual edits are valid under the emulator by construction, since any byte assignment yields a well-defined deterministic re-run. Off-manifold interventions are still valid for oracle scoring. The oracle measures a causal effect, not a likelihood. We ask what the output would be if u were v' . The bit-exact machine answers this exactly for any v' , on- or off-manifold. Where a method’s own claim is about typical behavior rather than causal dependence we fall back to the on-manifold donor and resampling variants and flag the choice in the Supplement.

The content-path gradient $\partial y / \partial u$ is the second way to build the oracle. We read it through the differentiable substrate. It is fast and dense, and it agrees with the intervention oracle on the content path, the register, color, and graphics-bit values that flow smoothly into the pixel. Yet, one important limit applies to the whole study. The straight-through gradient routes to the content path only. So a discrete index, a sprite’s screen position the VCS sets by the timing of a strobe write, has a naive gradient that vanishes under the hard forward map. We make this precise.

Proposition 1 (Zero naive gradient for discrete-index outputs). *Let an output y depend on a cause u only through a discrete index $n(u) = \lfloor g(u) \rfloor \in \mathbb{Z}$, the integer count of color clocks between a strobe write and the position-counter reset that the VCS uses to place a sprite, missile, or ball, with g continuous, piecewise linear, and non-constant (strictly monotone on each linear piece). Under the hard (bit-exact) forward map the rendered output is $y = h(n(u))$ for a fixed render function h . Then $n(\cdot)$ is locally constant wherever $g(u) \notin \mathbb{Z}$, so $\partial y / \partial u = 0$ at every such point, and y is discontinuous on the set $\{u : g(u) \in \mathbb{Z}\}$, which is measure zero precisely because g is*

strictly monotone on each piece, and there the gradient is undefined. Hence the naive gradient is zero wherever the derivative exists: a tiny change in u cannot move the sprite by part of a pixel. Any method that reduces to this naive gradient inherits the failure on position, index, and event outputs.

The proposition states only the hard-map limit; it does not say that no useful gradient exists. Two surrogate constructions of Paper 1 [4] are explicit exceptions and are *not* covered by it: the straight-through estimator, whose forward pass stays bit-exact to the hard machine at any finite temperature while routing a non-zero gradient along the content path [49, 50], and the bilinear sampler that interpolates the index boundary to recover a usable position gradient [51]. The companion paper constructs both and validates the bit-exactness this proposition assumes. We import that construction and prove the vanishing inline, so the present paper stands alone. We therefore avoid any blanket claim that “gradients are provably zero.” We use instead the scoped claim of Prop. 1: the *naive* gradient for discrete position/index outputs is zero under the hard forward map, and methods that reduce to it inherit that failure. Surrogate and bilinear-sampler gradients are exceptions.

We make the faithfulness audit as strict as possible. The position regime leaves one objection open, and we close it. A reader might object that the naive zero of Prop. 1 is just a fixable artifact. So we fix it: we turn on the Paper-1 bilinear sampler. The sampler interpolates the index boundary of Prop. 1 in the spatial-transformer manner. This makes the discrete position output a continuous function of the position RAM byte rather than a step. So the gradient methods receive a non-zero surrogate position gradient where the hard map gave them nothing [51]. We score both conditions, the naive forward and the sampler-on forward, against the same intervention oracle of Eq. (2). For each candidate cause the per-cause attribution is correlated, by the raw Pearson ρ of Section 3.4, against the oracle’s true $|\Delta \text{position-pixel}|$. So the two conditions differ only in the surrogate the method is handed, not in how it is judged. Alongside faithfulness we record a second, deliberately stricter measure. The *semantic recovery* of a method asks whether the method emits a Tier-3 game-concept name on its own, the word “ball” or “missile” attached to the byte it points at, rather than a cause that an external label only names later. An attribution map, which returns weights over causes and never a name, scores semantic recovery 0 by construction. Faithfulness and semantic recovery are two independent axes. Faithfulness asks whether the method names the right cause. Semantic recovery asks whether it names any meaning. The keystone experiment of Section 7 reports both for the naive and sampler-on conditions on the position regime.

Hence, for content outputs the gradient is a legitimate companion oracle, while for position, index, and event outputs the intervention oracle of Eq. (2) is the sole ground truth and every gradient-based method is scored as a method under test. This position regime is where the sharpest separations in our results appear. Where the two overlap we report their correlation and flag, rather than hide, the points at which they disagree.

3.3 Tiered ground truth and how we obtain it

The ground truth comes in three tiers. The tiers are not just bookkeeping; they support the paper’s central argument. Tier 1 (T1) is the exact causal map of Eq. (2), derived internally from the substrate and available for all 64 games. Tier 2 (T2) is the hardware roles and the read/write data-flow graph, also exact and internal across all 64 games. T1 and T2 describe what the machine does internally, and need no external label.

Tier 3 (T3) is different in kind. This difference matters for our argument. It attaches a game-concept name to a cell or feature, “the ball’s x position” or “the score.” These names cannot be computed from the substrate alone. They are facts about what the program *means*, which is separate from what it *computes*. We import T3 from two public sources and make it causal. AtariARI provides labels for 22 games, read from commented disassemblies, with raw RAM coordinates not aligned to the rendered position [6]; OCArari covers more than 40 games and, in its RAM mode, applies the render-time offsets that fix that misalignment [7]. We import both as candidates and verify each causally by setting the byte, re-rendering, and confirming on the bit-exact framebuffer that the named object moves to the predicted place. This upgrades a correlational label to a verified-causal one and auto-corrects the offsets; further labels we discover by RAM-to-framebuffer correlation and intervention. The headline core set is six games, Pong, Breakout, Space Invaders, Seaquest, Ms. Pac-Man, and Q*bert, each carrying both sources, all in the Paper-1 bit-exact roster. Together they span paddle-and-ball, shooter-with-shields, resource-meter, multi-agent-pursuit, and discrete-grid logic. The remaining games form a breadth set scored on T1 and T2 only.

We are careful about what this gives us. All T3 labels are external and only verified, never discovered from scratch. So a method that matches a supplied concept has aligned to a given meaning, not recovered it. This is exactly why a probe can show that information is *present* without showing it is *used* [8]. Our intervention test asks whether the byte *causes* the named object to move, so it is strictly stronger than decoding it. Recovering the program’s documentation and design from the bare machine is a separate problem, the subject of the companion paper.

3.4 The $F \wedge S \wedge M$ correctness triad

A faithful attribution is not yet an understanding. To measure understanding, we need an operational criterion. We adopt one in the spirit of Marr’s levels of analysis [47] and of Jonas and Kording’s warning that finding structure is not the same as explaining it [2]: each axis is first a continuous score in $[0, 1]$, and the boolean “right” is a thresholded reading of those scores. We score an explanation $\hat{E}$ of how an output y arises by three numbers, $F_{\text{score}}(\hat{E})$, $S_{\text{score}}(\hat{E})$, $M_{\text{score}}(\hat{E}) \in [0, 1]$, and call it *right* when all three clear their thresholds,

$$\text{right}(\hat{E}) = \mathbf{1}[F_{\text{score}} \geq \tau_F \wedge S_{\text{score}} \geq \tau_S \wedge M_{\text{score}} \geq \tau_M], \quad (3)$$

with $\tau_F = \tau_S = \tau_M = 0.5$ throughout. Results report the continuous scores; the predicate is a derived summary, not the primary quantity. All three are defined against the oracle, not against any method under test, which keeps the criterion non-circular.

Faithfulness asks whether the claimed causes are the *true* causes, scored against the intervention oracle of Eq. (2). The per-record metric depends on the explanation’s output type, so F_{score} is a single number assembled from three type-matched primitives. The committed score is the bare effect-agreement primitive, not a rescaled one. For a real-valued attribution map we report the Pearson (or Spearman, for ranked outputs) correlation $\rho(\hat{E}, \{\Delta y(u)\})$ between the attribution and the true causal effects, scored *raw*. A method whose attribution carries no causal signal therefore lands at $\rho = 0$, not at the 0.5 of an affine $\frac{1}{2}(1 + \rho)$ remap. This is what makes the position regime informative. Under Prop. 1 the naive gradient is identically zero there, and its correlation with the oracle is 0. The reported faithfulness is therefore 0.00 and not a 0.5 floor. The two committed estimators differ only in how the raw correlation is summarized across the six core games. Phase B reports the raw correlation directly (e.g. occlusion at $\rho = 0.482$), so a negative correlation is recorded as a negative number. In Phase A’s single-unit and tuning analyses the per-game correlations can be negative, so we clip each per-game correlation at 0 and then average across games (e.g. the tuning analysis A3, lowest in Phase A at 0.116, the clip-at-zero mean of six per-game correlations that include negatives). For the same attribution map we also report, as quality measures, deletion/insertion area-under-curve and precision at k against the true causal top- k , $P@k = |\text{top}_k(\hat{E}) \cap \text{top}_k(\{\Delta y(u)\})|/k$. For a graph-valued explanation (Phase A connectomics, Phase C circuit discovery) we use the edge F_1 against the true data-flow graph, the harmonic mean of precision and recall over recovered edges. For Phase C patching and mediation we use the causal-effect-agreement correlation against the true $\{\Delta y(u)\}$. The final F_{score} is the per-record primary metric averaged first within a game over its output records, then with equal weight across the six core games, so no game-rich method dominates; confidence intervals are 95% bootstrap over games (Section 3.5). The three Phases differ only in which primitive is primary: Phase A uses graph F_1 for the data-flow methods and the clip-at-zero mean correlation for the single-unit and tuning analyses, Phase B uses the raw effect-agreement correlation ρ with deletion/insertion AUC and $P@k$ as the attribution map’s quality measures, and Phase C uses the causal-effect-agreement correlation for patching and mediation and graph F_1 for circuit discovery. The leaderboard score is therefore a single scalar per method, but it is multi-metric in origin, and the per-method choice is recorded in the Supplement.

Sufficiency asks whether the explanation could regenerate the behavior under interventions it has not seen. We treat S_{score} as a held-out predictive estimator: $\hat{E}$ is fit (or its causal sites read off) on a calibration set of interventions and then asked to predict the true output y' under a disjoint test set of held-out $\text{do}(u)$ edits, with the prediction compared to the bit-exact re-run. The score is the fraction of held-out interventions whose predicted y' matches the true y' within a fixed output-type tolerance (exact for discrete indices and events, a small ε band for continuous scores), so $S_{\text{score}} \in [0, 1]$ and a value near zero means the explanation predicts no better than the unperturbed output. We deliberately do not define it through interchange interventions or scrubbing. Those are themselves mechanistic-interpretability methods under test in Phase C, so using them would make the criterion circular. Only Phase B and Phase C report S , and its bootstrap intervals follow Section 3.5.

Minimality asks whether the explanation is parsimonious and lives at the algorithmic level of the machine’s own module decomposition. We define $M_{\text{score}} = |U^*|/|U_{\hat{E}}| \in (0, 1]$, the ratio of the size of the true minimal cause set U^* , the smallest set of registers, RAM cells, opcodes, or program variables whose intervention reproduces y , to the size $|U_{\hat{E}}|$ of the cause set the explanation actually names. An explanation that names exactly the minimal set scores 1; one that records the whole machine state pays for every spurious cell. The target level is the algorithmic unit, the register, RAM cell, opcode, module, or program variable, not the pixel. Whole-state recording names all 128 RAM cells, where the true cause set averages a handful. It therefore scores $M_{\text{score}} = 0.07$ across the six core games. Its faithfulness is perfect because it withholds nothing, but it is maximally non-minimal.

Two reporting axes carry the headline comparison, faithfulness against the ground truth and a plausibility proxy. The plausibility proxy is a documented, tradition-level value, not a human-subjects measurement, and we report it as such. It is assigned per method tradition, not per method instance, from the published evidence on how convincing each tradition’s explanations are taken to be, before any faithfulness score on our machine is computed, so the proxy cannot be tuned to the result. The per-tradition assignment, its rubric, and its sources are tabulated in the Supplement (Sec. S7). There a sensitivity analysis perturbs the proxy under additive, rank-preserving, rank-jittering, and leave-one-tradition-out schemes. It finds the faithfulness-versus-proxy rank correlation stays negative throughout (from -0.79 to -0.40), so the headline contrast survives every alternative proxy assignment we tried. The proxy lets us find explanations with high plausibility but low faithfulness: explanations that convince but are wrong. F is satisfied by the wiring; S and M must be satisfied by the *computation*. Separating them lets us measure, later, what faithfulness alone leaves out.

3.5 Methods under test, games, and aggregation

We apply three traditions to the machine and score each by the triad of Eq. (3). The neuroscience battery (Phase A) runs connectomics and data-flow recovery, single-unit lesions, tuning curves, pairwise and global correlation structure, pooled-activity spectra, Granger causality, dimensionality reduction, and whole-state recording, scoring each against the known mechanism to quantify Kording’s lesson at the low-faithfulness baseline [2]. The attribution and XAI battery (Phase B) runs gradient saliency and its variants, integrated gradients, occlusion, perturbation masks, randomized masking, LIME, and Shapley-style attribution [10, 11, 16–20]; Grad-CAM and its variants and attention rollout [14, 15, 52] require neural layers or a learned policy and are recorded as not applicable. The mechanistic-interpretability battery (Phase C) runs activation patching and causal mediation, interchange interventions and distributed alignment search, attribution and path patching, automatic circuit discovery, sparse autoencoders, causal scrubbing, linear probing with control tasks, and the logit and tuned lens [8, 22–31, 53].

All faithfulness scoring runs on the six core games of Section 3.3, where the oracle’s re-runs are exact. Cell- and pixel-level scores, correlation, deletion and insertion area-under-curve, and precision at k , are taken against T1 and need no T3; only object- and concept-level scores draw on the external T3 labels, and we flag every such score as

alignment-to-given, not discovery. We aggregate the committed per-game records into one row per method on the shared faithfulness-versus-plausibility axes of the Results leaderboard. The headline contrast is the position regime, where the naive gradient is zero under the hard forward map by Prop. 1. There, causal and intervention methods average a faithfulness of 0.41 against 0.07 for gradient and correlational methods, a gap of 0.34. The per-method extreme places an exact causal method, activation patching, at faithfulness 1.000 and the field’s default saliency at 0.000 on all six games, the popular method nonetheless carrying the higher plausibility proxy (0.9 versus 0.5). We measure the semantic gap relative to these numbers.

4 The neuroscience battery, quantified

We begin with a known concern in the field. Jonas and Kording asked whether a neuroscientist with the standard toolkit could understand a microprocessor. They found that the toolkit yields rich structure but misses the mechanism [2]. They could show the failure but not measure it. On the 6502 they had no causal ground truth to score each method against. We supply that missing number. We run the same battery on the bit-exact, fully known VCS, where exact intervention recovers the true causal role of every register and RAM cell (Section 3). This makes Kording’s qualitative lesson quantitative. Across the eight analyses A1–A8, faithfulness against the oracle ranges from 0.116 to 1.0 and averages 0.49 (95% bootstrap CI ± 0.24) [36]. The oracle-as-method positive control sits at 1.0 by construction (Fig. 2). The structure these methods report is real. But structure is not the same as cause. The score measures the difference.

The dependency structure is largely recoverable. A1 reconstructs the dependency graph over state variables, a connectome of read/write edges, and scores it against the true data-flow graph at a faithfulness of 0.41 [36]: about half its edges are spurious. The single-unit lesion A2 fares best. It clamps each RAM cell in turn and ranks it by the behavioral damage it does. The resulting importance map has a Spearman correlation of 0.99 with the unit’s true causal role (mean 0.987 over the six core games) [36]. Lesioning is causal, which is why A2 is the one classical method that scores near the top.

Yet a near-perfect rank order still misses interactions. A single-unit lesion changes one cell at a time, so it is blind to causes that act only together. On Space Invaders the per-unit ranking misses 40% of the interactions the oracle records. One super-additive pair swings behavior by up to 1044 screen-pixels beyond the sum of its parts [36]. The ranking of which units matter is correct, but the model of how they combine is missing. So even the strongest classical method recovers a ranking, not the computation. An effect map does not show the algorithm behind it.

Tuning curves report structure that does not match the true cause. A3 fits per-unit tuning to a game variable, the workhorse of systems neuroscience. On the VCS it actively misleads. Its *spurious-tuning rate* is the fraction of strongly-tuned units whose tuning does not match their true causal role. That rate is 1.0 on Pong, and its faithfulness falls to 0.12, the lowest in the battery [36]. Tuning is correlational. On

a machine where many cells co-vary with the beam clock, apparent tuning is cheap. This is the same present-versus-used distinction we discuss later for linear probing.

Correlation structure and field potentials are real, but they are mostly side effects, not causes. A4, the spike-word analysis of pairwise and global correlation, reproduces the familiar weak-pairwise/strong-global signature. Yet against the true coupling its faithfulness is only 0.29, with sufficiency near zero ($S = 0.018$) [36]. A5 is the local-field-potential analogue that pools regional activity into power spectra. It scores 0.40, and much of that variance is the known frame and scanline clocks [36]. Both methods detect the machine’s activity but not its computation.

Granger causality treats “happens earlier” as “is the cause.” A6 infers directed CPU/TIA/RIOT edges and scores them against the true data-flow. On Pong its faithfulness is 0.0, with false-edge and missed-edge rates both at 1.0 [36]. The VCS is clock-locked, so nearly everything is predictable from everything else one cycle earlier. Equating precedence with causation therefore infers spurious edges everywhere. Neither a longer lag nor a stricter threshold fixes this. Precedence and causation differ here, and only the intervention oracle separates them.

Dimensionality reduction scores in the middle, with NMF and PCA tied. A7 decomposes the state tensor into latent components and matches them to the known clock, read/write, and vsync signals. Non-negative matrix factorization and PCA both match about 60 percent of the known components on average (at best 80 percent on a single game), so neither factorization is the more faithful at this analogue [36]. NMF leads only on the secondary recon-error/FSM composite, where its non-negativity better fits the additive register basis. This makes A7 the second-highest score, behind only the causal lesion. It still leaves 40 percent of the components unmatched.

A complete record is still not an explanation. A8 records the full RAM and register state, faithful and sufficient by definition ($F = 1.0$, $S = 1.0$). Yet its minimality is only 0.07 [36]. At the algorithmic level, a recording of every cell explains nothing. Faithfulness and sufficiency without minimality give a log file, not an account of the computation. This is the first sign that the three axes are independent. The semantic-gap result makes that point central later.

Taken together, the battery confirms the expected result and puts a number on it. The advance over Jonas and Kording is that we grade the *architectural* register-transfer level with a *quantitative* F against an exact oracle, rather than giving a qualitative verdict on the transistor. The contrast is clear, because the causal operators we apply to the same machine in Sections 5–6 reach the ceiling this battery cannot. So the limitation lies in the analysis methods, not in the machine. We mean this strictly as calibration. Even the battery’s best members recover a ranking, a graph, or a record, never the algorithm the program runs. This remaining gap is what the paper measures.

5 Results B — attributing a VCS output to its causes

Attribution returns a map of which inputs mattered. We apply these methods to the VCS to ask a question a real network cannot answer: is the map true? Every attribution is scored against the same intervention oracle (Section 3.2) that defines ground truth.

The neuroscience battery on a known machine — scored against ground truth

Jonas & Kording's battery, quantified: classical analyses find rich structure yet score low in faithfulness to the true mechanism.




Fig. 2 The Kording neuroscience battery scores well below the intervention oracle. Faithfulness of each classical analysis A1–A8 against the intervention oracle (Section 3.2); the battery averages 0.49, and only the causal single-unit lesion (A2) and dimensionality reduction (A7) sit clearly above chance, while the oracle-as-method positive control reaches 1.0 (dashed reference). Per-metric meanings: A1 = data-flow-graph F_1 vs oracle; A2 = lesion-importance ρ vs true role; A3 = 1–spurious-tuning rate; A4 = coupling correlation vs true coupling; A5 = 1–clock-explained variance; A6 = 1–false-edge rate vs true data-flow; A7 = NMF matched-component fraction; A8 = minimality M . Bars are coloured by method tradition; whiskers are asymmetric 95% bootstrap CIs (leaderboard.ci.csv), present on all eight Phase-A bars. Panel (b) is a key-results table isolating the mechanism behind the low scores: A6 Granger causality reaches $F' = 0$ on Pong (temporal precedence read as causation on a clock-locked machine); A3 tuning shows a spurious-tuning rate of 1.0 on Pong; A2 attains $\rho = 0.99$ yet misses 40% of interactions, one super-additive pair swinging up to 1044 pixels; and A7 NMF and PCA tie at 0.60 on matched components, both below two-thirds. The figure is the calibration baseline that the causal contrast of Fig. 4 sharpens. Every plotted value reads from a committed per-game record; no number is re-run or estimated.

For each chosen output y , a pixel, a score, or a game event, we know the exact set of causes it should have named. We ran twelve attribution methods across six core games (Pong, Breakout, Space Invaders, Seaquest, Ms. Pac-Man, Q*bert) and an oracle-as-method positive control, placed on the shared faithfulness axis in Figure 3 (Phase-B half) and scored per method in Table 1.

The main result is a split between two output types, not a ranking of methods. A content output is the value of a register, a colour, or a graphics bit. It is a differentiable function of the state, so its gradient flows. A position output is where a sprite is drawn, set by strobe timing on a discrete index. It has provably zero naive gradient (Paper 1, and Sec. 3). Hence the whole gradient family hits this limit, and it splits the leaderboard into two clear groups.

On the content path the gradient is faithful, as expected. Vanilla saliency [10] puts most of its attribution on the true causal byte: on Q*bert its attribution of a content read correlates with the oracle at Pearson 0.999 with precision@3 of 1.0, on Ms. Pac-Man at 0.811 / 1.0, on Breakout at 0.750 / 1.0. Integrated Gradients [11], Grad×Input/DeepLIFT [54], SmoothGrad [12], and guided backpropagation [13] all

Attribution (Phase B) vs mechanistic interpretability (Phase C): faithfulness on the bit-exact VCS

The 22 Phase-B + Phase-C method rows (a subset of the 30 methods + the oracle = 31-row leaderboard; the cross-tradition scatter is Fig. 2).

On the POSITION/INDEX regime (discrete sprite-position outputs, naive gradient provably zero): the whole gradient family collapses toward chance while intervention methods hold. Causal/intervention bucket 0.412 (95% CI ± 0.390 , n=4) vs gradient/correlational 0.068 (95% CI ± 0.070 , n=9) — a 0.344 faithfulness gap. Activation patching = 1.000 (oracle ceiling) vs vanilla saliency = 0.000 (naive-gradient floor).




Fig. 3 Attribution (Phase B) and mechanistic interpretability (Phase C) scored against the same intervention oracle (Section 3.2). The Phase-B half (a) splits each method’s faithfulness by output regime: on a *content* output the gradient family identifies the true cause, but on a *position*/index output the entire gradient family *collapses to 0.00*. The intervention methods still work, because each probe is a real re-run of the bit-exact program. Note that the popular gradient methods, which fail on position, have the higher plausibility proxy, the danger zone of Sec. 7 in numbers. The Phase-C half (b) is discussed in Sec. 6. The figure shows the 22 Phase-B and Phase-C method rows, a subset of the 30 methods plus the oracle positive control (31 rows) of the full leaderboard. Uncertainty bands are 95% bootstrap CIs: per method in Phase C, and at the family level for the Phase-B position-regime means (causal 0.412 [0.04–0.59] vs. gradient 0.068 [0.01–0.13]). Vector figure; values read from the committed cross-tradition leaderboard.

reproduce this, because the output we explain is a one-hot read whose Jacobian $\partial y / \partial u$ is a constant pointing at the read index. Smoothing does not change the score ($\Delta \text{corr} \approx 0$), and guided backprop reduces to vanilla saliency. On this path all methods give the same answer.

The gradient family fails on the position output. Its attribution goes to zero. The naive gradient is identically zero on every game, so saliency, Grad×Input, SmoothGrad, guided backprop, and Integrated Gradients all collapse to a faithfulness of 0.0 in the position regime — the “Fail / zero on index outputs” prediction of Sec. 3, now a measured number. Averaged over all output regimes the cross-phase gradient family scores only 0.294 faithfulness against the oracle (Fig. 3; Phase-B gradient methods alone 0.298); on the position regime alone, 0.068. Returning an empty map for half of a game’s outputs is not partial success. It is a structural failure, exactly where the discrete game logic is.

Two further results make this point stronger. The baseline choice is the perennial worry about Integrated Gradients. Here it changes magnitude, not faithfulness. The $(x - x_0)$ factor moves $\max |\text{IG}|$, yet the ranking and the oracle correlation hold, because the constant one-hot Jacobian keeps the mass on the same byte. Only the degenerate

baseline equal to the content byte breaks it. Expected Gradients [55] averages IG over a distribution of recorded VCS states. It removes that degeneracy and is provably stable across reference draws (mean pairwise cosine ≈ 1). Yet, on the content path it scores the same faithfulness as a well-chosen single-baseline IG at $M \times$ the compute, and it still vanishes on position. The more careful method gives stability but not accuracy. Its all-regime faithfulness of 0.034 is the lowest in the study.

Guided backprop and vanilla saliency fail the program-randomization sanity check of Adebayo et al. [34] on the system itself. We scramble the ROM so the program boots to a completely different state (~ 69 – 127 of 128 RAM bytes change, lowest on Seaquest), and the content-path saliency map stays byte-identical (cosine ≈ 1.0 on all six games). The Jacobian $\partial y / \partial u = e_{\text{idx}}$ is a constant one-hot, independent of what the program computes. So the map points at the read index no matter which program produced the output. The saliency is therefore provably model-invariant. A faithful explanation must change when the model changes. This one does not. The intervention oracle, by contrast, distinguishes the two programs at once. The map identifies the right cause but does not depend on the model that produced it.

When we replace the gradient with a real intervention, the position failure disappears. Occlusion [16] sets a candidate cause and re-runs the bit-exact program, a coarse version of the oracle. On the Pong position output, where every gradient method scores 0.0, it recovers the true causal bytes at correlation 0.837, its all-regime faithfulness of 0.482 the highest in Phase B. The black-box perturb-and-re-run methods follow the same logic and result: RISE [18] (0.375), LIME [19] (0.382), and KernelSHAP / Shapley sampling [20, 56] (0.446) each draw masked coalitions and average over the genuine responses, so each works on position outputs where the gradient is dead, and each reports its own completeness ($\Sigma \phi = y(\text{intact}) - y(\text{all-masked})$ for KernelSHAP, the surrogate R^2 for LIME).

The two minimal-mask methods meet Atrey et al.’s [3] objection that perturbations stray off-manifold. Extremal perturbation [17, 57] optimises an area-bounded mask of real *do*(occlude) re-runs, and the on-distribution counterfactual [58] substitutes candidate cells toward a genuine alternative VCS state, the RAM of another frame of the same ROM. Both stay on the data manifold by construction and both work on position outputs (position-regime faithfulness 0.202 and 0.139, all-regime 0.346 and 0.350). As such, the cross-phase intervention family averages 0.681 faithfulness [95% CI 0.59–0.78, $n = 11$] against the oracle, more than double the cross-phase gradient family’s 0.294 [0.20–0.38, $n = 14$] (Fig. 3). Restricted to the three Phase-B intervention methods alone the mean is 0.393 against the Phase-B gradient family’s 0.298. A perturbation is faithful only when it is a valid intervention.

Some popular methods return no answer at all, because the VCS has no structure for them to use. Grad-CAM and Grad-CAM++ [14, 15] need a final convolutional layer to weight; attention rollout [52] needs attention matrices; VIPER [59] needs a learned policy to distil. The VCS is a 6502 CPU with a TIA and a RIOT executing a fixed ROM. Across all six games the count of each required substrate is exactly zero, against ≥ 16 genuine candidate causes per game for the applicable methods. We record this as a measured finding, **applies = false**, not a poor score: the absence is structural and the same in every game, and re-targeting Grad-CAM’s pooling onto raw register

Table 1 Phase-B attribution scored against the intervention oracle (Section 3.2) (six core games; oracle-as-method positive control = 1.0 on every non-degenerate column). *Faith* is the all-regime faithfulness against ground truth and *Faith (pos)* the position/index regime alone, where the naive gradient is provably zero; *Plaus* is the documented method-tradition plausibility proxy (Sec. 3), *not* a human measurement. *Note that* the entire gradient family scores 0.000 on position while the intervention family does not, and that the popular gradient methods carry the *higher* plausibility tag — the danger zone, measured. The two family-mean rows are the *cross-phase* tradition families ($n = 14$ gradient/correlational, $n = 11$ causal/intervention across Phases A–C), not the twelve Phase-B rows above; the Phase-B-only gradient and intervention means are 0.298 and 0.393. Values are read from the committed cross-tradition leaderboard.

Method	Tradition	Faith	Faith (pos)	Plaus (proxy)
Occlusion [16]	intervention	0.482	0.306	0.55
KernelSHAP [20]	gradient/black-box	0.446	0.279	0.90
LIME [19]	gradient/black-box	0.382	0.167	0.90
RISE [18]	gradient/black-box	0.375	0.169	0.90
Grad×Input / DeepLIFT [54]	gradient	0.367	0.000	0.90
On-distribution counterfactual [58]	intervention	0.350	0.139	0.55
Extremal perturbation [57]	intervention	0.346	0.202	0.55
Integrated Gradients [11]	gradient	0.279	0.000	0.90
Guided backprop [13]	gradient	0.267	0.000	0.90
SmoothGrad [12]	gradient	0.267	0.000	0.90
Vanilla saliency [10]	gradient	0.267	0.000	0.90
Expected Gradients [55]	gradient	0.034	0.000	0.90
<i>Cross-phase family</i> — gradient/correlational		0.294	0.068	—
intervention		0.681	0.412	—
Grad-CAM / attention / VIPER [14, 52, 59]	N/A	does not apply (no conv / attention / policy)		

state would only be the content-path gradient renamed. Several well-known methods produce no result on a real artifact. This absence is itself a finding of the audit.

Across the twelve methods, faithfulness and plausibility move in opposite directions: methods rank high on plausibility but low on faithfulness. The clearest pair makes it concrete. On the position regime the causal method activation patching (Phase C) reaches faithfulness 1.000, the oracle ceiling, on all six games, while vanilla saliency, a widely used baseline, collapses to 0.000 on all six. This is a per-method gap of 1.000, bit-exact on every game. Yet, on the plausibility proxy of Sec. 3 the popular method has the higher score (0.9 versus 0.5): the empty map has a higher plausibility score than the exactly correct one. Low faithfulness with high plausibility is the danger zone. On the VCS we can measure it directly, not just warn about it.

We claim only that these maps are faithful, nothing more. A correct occlusion map names the true causal bytes. It does not tell us what those bytes mean, that a cell is the ball’s x or that a routine is the bounce. Object-level semantics enter only through the external T3 labels (AtariARI, OCArari; Sec. 3), and even there we verify a supplied label by intervention rather than discover it. Whether a faithful attribution amounts to understanding the computation we take up in Sec. 7. Here we have established only that faithfulness is achievable, and that the most popular methods do not achieve it.

6 Mechanistic interpretability on a known circuit

Mechanistic interpretability makes the strongest claim in the field. It tries to explain how a system computes, not just where it attends. Its methods read internal activations as a circuit and try to recover that circuit. On a neural network the recovered circuit can only be checked against itself, because the true wiring is unknown. The VCS removes that limitation. Its data-flow is the program’s data-flow, known exactly from the disassembly. Every intermediate value is a RAM cell or a register we can read, clamp, and re-run bit-exactly. We therefore run ten mechanistic-interpretability methods on a circuit whose answer we already know. We score each against the same §4 oracle: the exact intervention table, the true read/write graph, and the $F \wedge S \wedge M$ triad of §3. This calibrates the toolkit against a *known* circuit inside a real, hand-authored artifact rather than a synthetic or compiled one. Figure 3b places the Phase-C methods on the shared faithfulness axis. The result splits cleanly. The causal methods are exact, but their exactness recovers the wiring of the circuit without recovering what the variables mean.

Exact recovery of the wiring

The main result is simple. Activation patching recovers the true causal-effect table exactly. Across all six core games the recovered per-site effect equals the exact intervention oracle to the last bit (`recovered_equals_exact`: true; maximum absolute deviation 0.0), with site precision and recall both 1.0. On this machine, clamping a RAM cell to a donor value and re-running the real ROM is literally the oracle’s own operation. So it cannot disagree with the oracle. Causal scrubbing inherits the same strength. The true data-flow hypothesis preserves behaviour under resample-ablation of everything outside it on all six games (preserved performance 1.0). A deliberately wrong hypothesis does not (0.52), and the two are discriminated on every game. Interchange interventions with distributed alignment search (DAS) verify the supplied alignment at interchange accuracy 1.0 against a mis-aligned control at 0.0. The logit lens reproduces the true intermediate value at fidelity 1.0 on every game. Four of the five causal methods sit at faithfulness 1.0, level with the oracle positive control (Fig. 3b). Where a method’s operation *is* an intervention, faithfulness is automatic on a fully observable machine. This is the first part of the problem, and it is solved here.

Approximate and search-based recovery

The causal account weakens as soon as a method replaces exact re-runs with an approximation or a search. Attribution patching, the gradient-linear surrogate for activation patching, correlates with the exact patch at only 0.50 and matches it inside the linear regime in just 0.19 of cases. Yet its recovered edges stay precise (edge precision 0.95, recall 0.70) because the linearisation is wrong about the size of an edge, not about whether the edge exists. Path patching, run in the IOI style as a single directed do-operation, recovers the directed edge set of the true routine at mean $F1 = 0.46$. On Q*bert it matches the empty direct circuit perfectly ($F1 = 1.0$). On Pong it recovers none of the six true edges ($F1 = 0.0$), because each receiver’s value is re-derived within the horizon, so no one-step path carries a surviving signal. ACDC discovers

the circuit automatically by resample-ablation edge pruning over a threshold sweep. It reaches mean $F1 = 0.45$ (ROC-AUC 0.70). A positive control re-runs the identical edge-restricted procedure under the exhaustive do-set and scores $F1 = 1.0$ on every game, so the deficit is a real measurement of single-shot recovery, not a broken scorer. These are the methods the field reaches for when the exact intervention is too expensive. On a machine where the exact intervention is cheap, these methods miss real structure. They are precise about what they find but silent about dependencies re-derived through other cells. The limit is recall, not precision. The gap to the positive control is the quantified cost of approximation against a circuit that is known.

Three separations from the oracle

This subsection presents the central result. ACDC on Breakout recovers the data-flow circuit at perfect $F1$ against the true routine, and the recovered graph is minimal at the algorithmic level ($F = 1.0$, $M = 1.0$). Yet that same graph, scrubbed and asked to regenerate the behaviour under held-out interventions, preserves only $S = 0.44$ of it (mean scrub-preserved performance 0.32 across games). A circuit can have the correct graph and still not reproduce the computed function. The missing object is not a denser or cleaner graph of the same type. Activation patching has already recovered the effect table exactly, with no tuning budget left to spend. A table of per-site effects is still not the bounce algorithm. One method is exact and the other recovers a perfect graph. Even so, neither reproduces the behaviour. A “methods underperform” explanation cannot account for this separation, because no method underperformed here.

The dictionary methods show the same gap from the opposite direction. A sparse autoencoder trained on the recorded RAM trajectory reconstructs the state almost perfectly (fraction of variance explained > 0.99 on Pong) and matches every verified game-variable cell it was asked to (matched fraction 1.0). By the metric the field uses to claim a feature “is” a concept, the SAE succeeds. By the oracle, it does not: ablating the matched Pong feature and re-running the real ROM gives causal faithfulness $F = 0.041$, with $S = -0.336$. When we remove the matched feature, the computation it is supposed to name does not change. Across the six games the SAE’s mean leaderboard faithfulness is 0.19 against a method-tradition plausibility proxy of 0.75, squarely in the §3 danger zone. The NMF/PCA dictionary contrast repeats the pattern at lower resolution (matched-component fractions up to 0.75 beside near-zero causal use). This separation is the reverse of the first. The circuit was causal but behaviourally incomplete; the feature is perfectly named but causally inert. The two failures point in opposite directions.

Linear probing shows a third version of the same gap, the well-known probing trap. With control tasks (Hewitt & Liang’s selectivity), probes decode game-concept cells well above chance (mean probe accuracy 0.89, mean selectivity 0.16), confirming that the information is present in the state. Yet the program does not necessarily use that information. Cross-referenced against the exact intervention oracle, the oracle flags 25 concept cells across the core games that the program never causally uses within the window. Of those, 6 are also decodable above chance, so a probe would have nominated them as encoded concepts. Breakout’s score cell 84 sits in the larger set: it is present in the state and linearly associated with the displayed score, yet provably

inert under intervention within the window. The oracle is strictly stronger than the probe. The probe asks whether a value can be read out. The oracle asks whether changing that value changes the output. Probing earns a high plausibility proxy (0.85) on a faithfulness of 0.16, the second-worst entry in the danger zone (Fig. 3b). We state the caveat directly: these labels are imported externally (AtariARI, OCArari) and only *verified* causally, never *discovered*. So DAS’s alignment accuracy of 1.0 is guaranteed in advance, because it only confirms a variable we supplied. The probe tells us a concept is present. Only the intervention tells us the program uses it.

Exactness is necessary but not sufficient

Phase C is the clearest test of the paper’s claim. Run against a known real circuit, the mechanistic toolkit splits into two groups. The first is a causal core that is faithful by construction: activation patching, causal scrubbing, DAS, and the logit lens, all at 1.0 (phase mean faithfulness 0.63, the highest of the three phases). The second is a recovery layer that is either precise but partial (attribution and path patching, ACDC) or present but inert (SAEs, dictionaries, probing). The oracle, not our judgement, decides the three separations. They appear exactly where the field’s success metrics report success: graph-correct yet behaviour-incomplete (ACDC on Breakout $F=1.0/M=1.0$, $S=0.44$), named yet causally inert (SAE on Pong matched 1.0, $F=0.041$), present yet unused (cell 84). The faithfulness advantage that the causal core holds over the gradient methods is itself regime-dependent, and we state that plainly. The gap is robust only on the position-output regime (0.34, 95% CI [0.02, 0.50], excluding zero). It collapses when the output is content rather than position or when the oracle-equivalent controls are excluded (0.12 and 0.04, both with confidence intervals crossing zero). The separations above survive that caveat because each is a within-method contradiction on a single game, not a cross-family mean. On a machine that gives faithfulness for free, faithfulness is necessary for understanding but not sufficient. The residual is the semantic gap: the wiring is recovered, but the meaning is not. §7 places this gap on the shared axes, and the Discussion treats it as the field’s open problem.

7 Across the traditions — causal structure without meaning

We have now run three traditions on one machine. This section is organized not by the ranking among them, but by a boundary they share. Every method, faithful or not, returns the system’s causal structure but not the meaning of its variables. Faithfulness is whether an explanation names the true causes. On the VCS we can measure it exactly. The intervention oracle of Section 3.2 re-runs the bit-exact program for every candidate cause. It returns ground truth that no real network can provide. That exact answer lets us do two things in sequence. First, we place all of Phases A–C on the shared faithfulness-versus-plausibility plane. We read off the leaderboard and its danger zone, the known result that faithful and convincing are not the same. Second, we go further. We make the unfaithful methods faithful by turning on the differentiable port’s bilinear sampler. Making every method faithful does not change the result, because faithfulness was never the dividing line. The harder part of the

problem is the variable’s meaning and the routine’s algorithm. This is a separate object, and no method on this fully known machine recovers it. We call the distance to it the *semantic gap*. The leaderboard is an intermediate result; the semantic gap is the main result. The correctness triad we apply throughout — **Faithful**, **Sufficient**, **Minimal** — is defined formally in Section 3.4; here we only read its scores.

The faithfulness leaderboard

The leaderboard separates cleanly along the faithfulness axis. We read it as an exploratory cross-tradition ranking, not a single commensurable score, because the traditions measure different objects. We report the ranking with its robustness. We aggregate the 257 committed per-game records into the canonical 30 interpretability methods plus the intervention oracle as a positive control, 31 rows in all [36] (Fig. 4). The causal and intervention methods average a faithfulness of 0.681 against the oracle, while the gradient and correlational methods average 0.294, a gap of 0.387 across all output regimes. The position regime is the set of discrete sprite-position outputs whose naive gradient is provably zero under the hard forward map (Section 5). On this regime the gap holds with the same sign: causal and intervention 0.412 versus gradient and correlational 0.068. The ordering is the one the method matrix predicted (Section 3), now measured. Activation patching, causal scrubbing, interchange interventions, and the logit lens sit at the oracle ceiling of 1.0; the single-unit lesion follows at 0.99; the gradient family and the classical correlational analyses (tuning curves, Granger causality, and spike-word correlations) score near zero. The clearest contrast is between two methods, and it holds bit-exactly on every game. On the position regime activation patching reaches faithfulness 1.000, the oracle ceiling, on all six core games, while vanilla saliency, a widely used attribution baseline, collapses to 0.000 on all six, a per-method gap of 1.000 [35]. Hence, on a system where the answer is known exactly, an everyday attribution tool scores near chance where a causal method scores near the ceiling.

The separation survives reweighting, and where it does not we say so. Reading the robustness record [36], the causal-versus-gradient gap stays at 0.387 under both equal weighting by method and equal weighting by game, and its 95% bootstrap interval excludes zero in each case ([0.32, 0.45]). Yet two qualifications must be stated alongside this result. When the oracle-equivalent intervention methods are removed, leaving only coarser approximations, the gap shrinks to 0.123 and its interval includes zero, so part of the leaderboard’s main contrast comes from methods close to the oracle operation itself. The gap is also regime-specific: on content outputs, where the naive gradient is valid, it falls to 0.045 with an interval spanning zero, while on the position regime it holds at 0.344 with an interval that excludes it. The conclusion that survives every reweighting is the narrow one: on the discrete position/index outputs, causal and intervention methods recover the true causes and gradient-and-correlational methods do not. This is the cross-tradition analogue of what neural-network benchmarks report when faithfulness is scored against a constructed ground truth [38, 41], with the difference that here the ground truth is an exact intervention on a real software artifact rather than a synthetic or pseudo-label.




Fig. 4 Every method, every tradition, on the shared faithfulness-versus-plausibility plane. The plotted set is the canonical 30 interpretability methods + the oracle = 31 rows, each placed by its faithfulness against the intervention oracle (Section 3.2; X, exact ground truth) and by its documented method-tradition *plausibility proxy* (Y; Section 3). The plausibility axis is a tradition-level proxy, not a human-subjects measurement. The dashed curve traces the offset between the faithful and the plausible clusters, and the oracle positive control anchors (1.0, 1.0), a ceiling no interpretability method reaches. Causal and intervention methods cluster at high faithfulness while the gradient and correlational methods cluster low yet *higher* on the proxy — the *danger zone* of high apparent plausibility at low faithfulness. The key findings and per-method values are in the in-figure table; every plotted value reads from the committed leaderboard, none re-run or estimated. Per-method uncertainty is reported in Fig. 6 and the supplement.

The danger zone

If the faithfulness axis matched how convincing a method looks, it would tell us nothing new. It does not match. Plotted against the plausibility proxy of Section 3, the points split into a safe diagonal and a dangerous corner. The methods the field uses first sit near zero faithfulness while having the highest plausibility-proxy values. Expected Gradients tops the danger list at plausibility 0.9 against faithfulness 0.034; linear probing follows at 0.85 against 0.162; A3 tuning at 0.8 against 0.116. This inversion is the key finding: the map provably empty on the position regime looks *more* convincing than the exactly correct map. We state the caveat plainly. The plausibility axis is a documented method-tradition proxy assigned per tradition before the faithfulness scores were read, not a human-subjects measurement, and we report it as such to avoid overclaiming the danger zone. Still, the faithfulness axis is exact, and the ranking it

inverts is what the field tends to trust, so the hazard is real even at the proxy’s coarse resolution. The danger zone is a quantitative result, not a speculation.

Making the unfaithful faithful

A reader could accept the danger zone and still dismiss it by blaming the gradient. The saliency methods score zero on the position regime because the hard forward map reads a sprite’s position through an integer index, so the gradient is zero almost everywhere (Section 5, Prop. 1). On this reading, the danger zone is a technical artifact a smoother forward pass would repair. Once repaired, the gradient methods would become faithful and the concern would disappear. The differentiable port lets us run exactly that repair and test the claim. We turn its bilinear sampler on, replacing the integer index read by a triangular-kernel occupancy whose sprite column is a continuous function of the position byte. This restores a non-zero gradient through the position computation without altering the forward output. The same surrogate that recovers the joystick gradient in Section 5 now flows through the position read.

The repair works on the games that have a moving, tracked sprite. On Pong, vanilla saliency, gradient $\times$ input, SmoothGrad, and expected gradients all rise from the provable zero to faithfulness 0.791 against the oracle; on Q*bert vanilla saliency, gradient $\times$ input, SmoothGrad, and integrated gradients rise to 0.529 (committed sampler record, `compare/out/sampler_faithfulness.csv`). The rise is not universal, and we report it as it is: integrated gradients with a zeros baseline still vanishes on Pong, the static-sprite games offer no position to restore, and on Space Invaders the restored gradient lands on a non-active cell, so it does not move. Where a sprite’s position is actually tracked, the previously empty saliency map becomes faithful. This rules out the gradient-blaming objection. The danger-zone methods are no longer unfaithful by construction.

But this does not give understanding. The restored gradient on Pong concentrates entirely on the single RAM byte that holds the paddle’s position, `ram[54]`, which the oracle independently confirms is a true cause. The method now finds the right cause and attributes nothing to the wrong ones. Yet its semantic-recovery score is 0, exactly as it was when the map was empty and exactly as it is for every other method in the study (same committed record). A faithful attribution to `ram[54]` names a byte the program reads. It does not say that the byte *is* the paddle, that reading it *is* tracking, or that the routine computes a bounce. The only semantic label in play is the imported annotation we use to *check* where the attribution landed, never one the method *produced*. We could repair faithfulness, but not meaning. This result sets up the rest of the section: the gap we measure next is not the unfaithful methods’ deficit, because making them faithful does not close it (Fig. 5).

Where the leaderboard stops — the semantic gap

Here the paper turns. Take the best case for the audit: a method that scores 1.0 against the oracle. We ask whether it gives an understanding of what the program does. It has not, and we are careful about what that claim is and is not. We do *not* claim that no method can recover meaning; that is the unbounded statement, the subject of the

Keystone: the bilinear sampler makes gradient attribution faithful on the position regime — yet recovers no semantics

On the discrete position/index regime the naive gradient is provably zero (Prop. prop:zero), so every gradient method scores faithfulness 0. Turning Paper-1's bilinear index-boundary sampler ON (same surrogate as the SI joystick tool) restores a real position gradient and the methods become FAITHFUL — up to 0.791 (Vanilla saliency on pong). Yet semantic_recovery stays 0 for all 30 (method $\times$ game) records in BOTH conditions: an attribution map names no game concept on its own. Faithfulness is repairable; the semantic gap is not — the danger zone is not a fixable artifact of the vanishing gradient.




Fig. 5 Faithful, still no meaning: turning the sampler on lifts faithfulness off the floor and leaves semantic recovery at zero. On the position regime the naive gradient is provably zero (Prop. 1), so the saliency methods score faithfulness 0 (left bars). Turning the differentiable port's bilinear sampler on restores the gradient through the position read and lifts faithfulness to 0.791 on Pong and 0.529 on Q*bert for the methods with a moving tracked sprite (centre bars), with the restored gradient concentrating on the true position byte (`ram[54]` on Pong, oracle-confirmed). The semantic-recovery score stays 0 in both conditions and for every method (right), because an attribution map names a cause but emits no game concept; the only semantic labels are imported (Section 3.2) and used to check localization, never produced. Note that the rise is reported only where a tracked sprite exists: integrated gradients with a zeros baseline vanishes on Pong, static-sprite games have no position to restore, and Space Invaders' restored gradient lands on a non-active cell. Values read from the committed sampler record; none re-run.

companion design-recovery work, which we have not run here. We claim the bounded thing the oracle can prove: even when a method recovers the true causal structure exactly, what it returns is the causal structure — a causal-effect map, a data-flow graph — and not the *semantics*, the variable's meaning or the routine's algorithm. The missing object is a different kind of thing. It is not just a poorer version of the causal structure. Following the triad of Section 3.4, we operationalize the gap through three separations, each adjudicated by the oracle: faithful-not-sufficient, named-not-causal, and decodable-not-used. They diagnose distinct semantic failure modes; they are not a complete theory of semantics. The three separations show the failures pointing in opposite directions (Fig. 6).

The first separation is graph-correct yet behavior-incomplete: high F and M , low S . Automatic circuit discovery (ACDC) [25] on Breakout recovers a data-flow graph that is, by the exact data-flow oracle, perfectly correct and perfectly parsimonious, $F = 1.0$ and $M = 1.0$ [36]. Yet the discovered circuit alone, with its non-circuit parents scrubbed, reproduces the clean candidate-cell readouts only 44% of the time, $S = 0.44$. A correct graph of which cells feed which other cells is not the function those cells compute. The faithful side shows the same. Activation patching [22] recovers the

exact per-site causal-effect table on every one of the six games [36], and a table of effects, however exact, is not the bounce algorithm that produced them. No tuning budget is left to spend here, because the effect table is *already* exact, so the residual to the algorithm cannot be a worse fit closed by a better hyperparameter. It is a missing object of a different type, and that is why the gap is not merely “methods underperform.”

The second separation is named yet causally inert, the reverse failure: high matched fraction, low causal F . A sparse autoencoder [28] trained on the recorded Pong state matches every one of its candidate game-concept variables, a feature-to-variable matched fraction of 1.0, a perfect semantic alignment by the field’s usual yardstick. By the oracle, those same matched features are nearly causally inert on Pong: faithfulness $F = 0.041$ and a *negative* sufficiency $S = -0.336$ [36]. A feature that looks like the ball’s x is not the same as the mechanism the program uses to compute with the ball’s x . The circuit was correct but incomplete; the dictionary is named but causally wrong. The gap runs in the opposite direction on the same axis. An account of methods simply scoring too low cannot explain this. Recent SAE evaluations report the same proxy-versus-use gap on neural language models [32]; the oracle here turns that observation into an exact causal-use test on a known system. We guard against reading semantic success into an aligned variable. The interchange-intervention alignment (DAS) [27] reports an aligned-variable accuracy of 1.0, but that 1.0 is true *by construction*, verifying a *supplied* variable rather than discovering one. A method that confirms a given concept is not a method that found the meaning.

The third separation is decodable yet unused, the present-versus-used trap, where the oracle is strictly stronger than the probe: high probe accuracy and selectivity, low intervention effect. Linear probing with control tasks [8] decodes game concepts from the Breakout state at a mean accuracy of 0.72 with positive selectivity over its control; the information is demonstrably *present*. Yet the exact oracle flags a score-related cell that is *present but unused* (cell 84): re-derived within the horizon, its value has no downstream effect on the output [36]. A concept can be linearly readable yet play no part in the computation, exactly Hewitt and Liang’s control-task warning; decodability is necessary, never sufficient. Probing answers whether the concept is present. The oracle answers whether it is used. Only the second tells us how the program works. Linear probing’s leaderboard faithfulness of 0.162 paired with a plausibility proxy of 0.85 is its place in the danger zone, and this cell is the mechanism behind that placement.

Read together, the leaderboard plane of Fig. 4 and the failure taxonomy of Fig. 6 make the same point from three sides, and the triad’s axes are not redundant: a method can be faithful and minimal yet not sufficient (ACDC on Breakout, $F = M = 1.0$, $S = 0.44$); it can match a named concept perfectly yet be causally inert (the SAE on Pong, matched fraction 1.0, $F = 0.041$); it can flag a concept the program never uses (cell 84, present but unused). The claim does not rest on any single low number; a lone low S could be a tuning artifact. It rests on two things. First, the gap appears where tuning cannot help: the effect table is already exact. Second, the failures run in opposite directions on the same axis. A “methods are not good enough yet” account cannot explain this two-sided result. We restate the honesty contract here, because

this is the easiest place to overread the results. We have *not* recovered the program’s meaning, algorithm, variables, or design from scratch. Every semantic label we used is external, imported from AtariARI [6] and OCArari [7]. The substrate only *causally verifies* a supplied label, it does not *discover* the concept, and the provenance record marks each label as imported rather than recovered. So the faithful methods recover the causal structure. The meaning is a separate thing, and we have measured the distance to it. Recovering that meaning is a different problem, and it is the companion paper’s, not this one’s.

A note on aggregation closes the section, because two committed records report the sparse autoencoder differently and a careful reader will see both. The leaderboard score [36] aggregates over the per-method records, which for the SAE gives faithfulness 0.195 with a 95% bootstrap-over-games interval of [0.00, 0.31] that brackets it; the figures plot this value. The per-game Pong record used in the named-not-causal separation above is a different, finer object, $F = 0.041$ on that one game. We keep the two distinct: the 0.195 is the leaderboard aggregate, the 0.041 is the single-game illustration. Throughout the paper, the leaderboard and its figures use the per-method aggregation as canonical, so text, table, and figure agree.

8 Discussion

We applied the interpretability toolkit to a machine whose ground truth is computable. This let us measure two things at once. On this system the true cause of any output is recoverable by exact intervention. The causal and intervention methods are faithful, while the field’s familiar gradient and correlational tools are not. The leaderboard separates them by faithfulness, the one axis we can score exactly (Fig. 4). A faithful causal account scores 1.0 where the field’s default scores 0.0, on a machine whose answer is known exactly. Yet faithfulness is the easier of the two measurements. Even a method at the oracle ceiling returns only the causal wiring: an effect map and a data-flow graph. It does not return the semantics: the meaning of a variable or the algorithm of a routine. We call this residual the *semantic gap*, and it is a number we can measure.

We score three traditions with one common oracle: a neuroscience battery, the attribution toolkit, and mechanistic interpretability. These traditions are rarely studied together, but the oracle places them on one scale. Jonas and Kording warned that an analyst armed with systems neuroscience could fail to understand even a micro-processor, because describing structure is not the same as explaining how it works [2]. We make that a measurement. The classical battery recovers rich structure yet scores low against the known mechanism. The gradient methods score zero on the discrete position regime, and only the causal operators reach the ceiling. Mechanistic interpretability has until now been calibrated against synthetic or compiled ground truth, transformers with a planted circuit by construction [39, 40], and ground-truth attribution benchmarks likewise plant the important features [37]. We add the missing case: a real, hand-written, deployed artifact that was never designed to be interpretable. Its ground truth is independent of the artifact, not built into it. So the leaderboard

calibrates these methods against ground truth on a system the field did not design to be interpretable.

The same oracle also reveals a limit: a high leaderboard score cannot show meaning. This limit is the subject of this paper. A perfect faithfulness score certifies that a method named the true causes. It does not certify that the method recovered what the program means. We observed this limit in three cases (Fig. 6). On Breakout a circuit is graph-correct and minimal ($F = M = 1.0$) yet reproduces behavior only 44% of the time ($S = 0.44$). On Pong a feature dictionary matches every named variable (matched fraction 1.0) while being nearly causally inert ($F = 0.041$, $S = -0.336$). And a concept is linearly decodable yet, by the oracle, never used in the computation. The gap appears in two opposite forms: correct-but-incomplete, and named-but-wrong. A simple “the methods are not good enough yet” story cannot produce both.

We tested the strongest version of that story and it failed. The faithfulness gap could be read as an artifact of one design choice. The discrete position read-out makes the naive gradient vanish by Prop. 1. So the gradient family scores a provable zero on the position regime, and the failure could be dismissed as a fixable technical artifact. We removed that artifact. Turning the port’s bilinear sampler on restores a usable gradient through the position index. On the games with a moving tracked sprite, the gradient methods become faithful. Their position-regime faithfulness rises above zero. On Pong it reaches 0.791 for vanilla saliency, grad×input, smoothgrad, and expected gradients. On Q*bert it reaches 0.529 for vanilla saliency, grad×input, smoothgrad, and integrated gradients (Fig. 5). The restored gradient concentrates on the position byte that drives the sprite, `ram[54]` on Pong; it finds the right cause. Still, the semantic recovery of every one of these methods stays 0, exactly as before and exactly like every other method in the study. We removed the faithfulness deficit, but the meaning still did not appear. Hence the shortfall is not a faithfulness deficit waiting for a better gradient. Even when the gradient is corrected, it identifies the cause but no meaning. So the meaning was never present in the wiring for a method to recover.

The improvement occurs only in specific cases, and that pattern is itself evidence. It is not a universal rise. It rises only on the games with a moving tracked sprite the gradient can flow to. Integrated gradients with a zeros baseline still vanishes on Pong. The static-sprite games show no rise, and so does the case where the restored gradient falls on a cell the program does not currently read. Where a position byte exists and moves, the sampler makes the gradient faithful. We gave the faithfulness explanation every case it could win. None of it produced a single semantic label (Fig. 5).

What understanding would require

Faithfulness stops short for a reason that Marr’s levels of analysis make clear [47]. A faithful attribution lives at the implementational level, where cells and registers carry the signal. Understanding lives one level up, at the algorithmic level, in what function the program computes. A correct data-flow graph is an implementational fact. The bounce algorithm it implements is an algorithmic fact, and the two are objects of a different type. Lazebnik made the same point: one can describe a system in complete mechanistic detail and still not understand it [48]. We turn that observation into a measured quantity. The distance from a faithful effect table to the algorithm is the

sufficiency gap $1 - S$, and on an already-exact effect table it cannot be closed by a better hyperparameter.

Software engineering gives the clearest version of this requirement. IEEE 610.12 defines software as programs plus their documentation and data, the specification, the design, the intent [9]. A faithful circuit recovers part of the program. To understand the system you must also recover its documentation: what each variable is for and what each routine should do. The sampler-on result tells us why no internal method can reach that object. The meaning of a variable is not a causal fact about the variable. It is a fact about the behavior the design serves, and that behavior is fixed outside the program code, in the manuals and the data the program was written against. A method that reads only the mechanism, however faithful, reads only causes. A cause becomes a meaning only against an external reference: the system’s documented behavior. No perturbation of the chip can reach that reference. In this study the only semantics present entered from outside: every T3 label was imported from AtariARI and OCAtari and then merely verified to be causally live, never discovered by a method. Hence the gap is not an artifact of weak methods but a property of the question. Faithfulness is necessary for understanding but not sufficient, because understanding needs an external behavioral reference that an internal method does not carry. We state clearly that this paper defines and measures the gap. Closing the gap means recovering the program’s documentation and design from the bare machine. That is the reverse-engineering goal, and the task of the companion design-recovery paper (Paper 4). That paper measures the distance to that goal rather than clearing it. Whether a recovered concept matches how the system is used is the companion behavior paper (Paper 3). The learned agent, where there is no disassembly to check against, is Paper 5.

Probing reads the same result from the data side. It shows that a concept is present in the state, but not that the concept is used. This is the control-task warning of Hewitt and Liang [8]. Our intervention oracle is strictly stronger, because it asks whether the byte causes the output, and it found a present-but-unused score cell the program never reads. Decodability is presence; use is an algorithmic fact. Understanding needs use, and only the oracle supplies it. At base, the three traditions are one toolkit under different names. Tuning curves match feature visualization, lesions match ablations, and Granger causality is the correlational predecessor of causal attribution. Each pair shares a failure mode, so a result measured on one transfers as a caution to the others. A benchmark built on an Atari chip thus speaks to systems neuroscience, explainable AI, mechanistic interpretability, and software verification at once.

Why a 1977 chip is a fair screen

We address the obvious objection. The VCS is deterministic, hand-coded, and small, while the systems these methods target are learned, distributed, stochastic, and large. We do not ask a score here to predict a transformer. The benchmark is a necessary-condition stress test, not a sufficiency predictor. Here a method has perfect ground truth, full observability, exact interventions, and no learning to complicate the result. If it still fails, that failure is strong evidence against any claim that the same method reliably recovers causal mechanisms in harder, less observable systems. Passing here is necessary but not sufficient; failing here is informative. The same logic suggests,

rather than proves, a reading of the semantic gap. If even the strongest causal account a fully observable machine permits stops short of meaning, the result suggests that faithfulness alone should not be treated as understanding in less observable neural systems, where the circuit itself is uncertain and no clean ground truth is on hand. We measure $1 - S = 0.56$ on an exactly recovered circuit on Breakout ($1 - 0.4444$; the ACDC family mean is $1 - 0.32 = 0.68$), and we do not extend that number to neural networks as a proven floor; the precise scope, and what this paper does not claim, are set out in the Supplementary box on what we do and do not measure (S6).

The chip also exhibits the properties that make real systems hard to interpret. There are aliased, polysemantic RAM cells reused across routines, a hand-coded analogue of superposition. There is computation by racing the electron beam, which is distributed and temporal. And there are registers that mean different things in different contexts. Figure 7 maps each VCS property to the neural-network difficulty it does and does not capture, so that a property we can prove on the VCS is one to suspect in neural networks. We state plainly what the chip does not capture: learned features, scale, and stochasticity are absent. The learned agent that adds all three back is Paper 5.

What interpretability needs

Three recommendations follow. The field should prefer causal over correlational evidence, because the methods it reaches for first carry the highest plausibility proxy and the lowest measured faithfulness, and a saliency map provably empty on the position regime can look more convincing than the exactly correct map. It should validate on known systems before trusting on unknown ones, since a method that cannot pass where the answer is computable cannot be audited where it is not. Most important, it should stop treating faithfulness as understanding. The open problem for explainable AI is not a better saliency map; it is the semantic gap. We confirmed this directly: we closed the faithfulness gap for the gradient methods, and the meaning still did not appear. The field does not need an instrument that reads the mechanism more sharply. It needs one that supplies the reference the mechanism does not contain: the system’s behavior, documented for software in its manuals and data. Faithfulness is the part we can now measure, and on causal methods we can achieve it. Meaning remains unsolved, because no internal method carries the external reference it needs.

We have measured a gap, not closed it, and several limits bound what we may conclude. We did not run semantic recovery, because Phase E, documentation and design recovery, is specified but not executed. Our claim is therefore the bounded one: on exact ground truth, faithfulness certifies the causal structure, and that structure alone is not the semantics. We do not make the unbounded claim that “no method can recover meaning,” which is Paper 4’s charge. The sufficiency and minimality axes are defined against the oracle, and the T3 semantic labels are external, imported from AtariARI [6] and OCArari [7] and only causally verified, never discovered. A single low sufficiency could be a tuning artifact. So the claim does not rest on any single number. It rests on three facts: the gap appears where tuning cannot reach it; it runs in both directions; and it survives the sampler-on test that makes the gradient methods faithful on games with a moving tracked sprite. We do not read that test as

universal. The rise occurs on Pong and Q*bert, but not on the static-sprite games, not on a zeros-baseline integrated gradient, and not where the restored gradient falls on a cell the program does not read. We report it as such, not as a law over all methods and games. The plausibility axis is a documented method-tradition proxy, not a human-subjects measurement, and we report it at that proxy’s coarse resolution. Our scores hold within Paper 1’s validated conformance window. Outside that horizon the substrate is not certified bit-exact, and we keep every scored claim inside it. The subject is one old machine with no learning, which is the necessary-condition framing once more: a stress test, not a forecast.

We end with the broader implication. Because the substrate is differentiable, it tells us which gradient explanations are trustworthy and where they vanish. The content path carries an honest gradient while the discrete index does not. So a saliency map over sprite position is empty for a structural reason, not by chance. This statement about gradients generalizes to any system where a discrete decision follows a continuous one. We believe a fully known, differentiable testbed is valuable for exactly this reason: it lets the community calibrate its instruments before deploying them where calibration is impossible. The hard case remains the learned agent, where there is no disassembly, no specification, and no engineer’s intent to recover, and where the semantic gap is likely to be larger. Closing this gap may take the next decade of interpretability research. The goal is to turn a faithful account of a learned system into an understanding of what it has learned. We have not solved this problem. We have defined and measured it, which is the necessary first step.

Data availability. The interpretability scores reported here come from the committed analysis records under `tools/xai_study/`. These records hold the per-method, per-game §R results and the aggregated cross-tradition leaderboard [35, 36]. Every quantitative claim in the Article reads from one of these committed artifacts. No number is re-run or estimated. The game cartridges used as the subject are third-party Atari 2600 ROMs that we do *not* redistribute. We follow the practice of the underlying emulator port [4]. We acquire the ROMs through the AutoROM retrieval procedure and report a SHA-256 hash for each ROM in the Supplementary Information. This lets others reproduce and verify the exact binaries independently. The recorded state trajectories used to train the sparse autoencoders and to run the patching and probing experiments are regenerated deterministically from those ROMs and the published action streams.

Code availability. The interpretability benchmark introduced here is available for review at <https://github.com/akmaier/UnderstandingVCS>. It comprises the scoring suite, the $F \wedge S \wedge M$ triad metrics, the intervention oracle, and the per-phase method implementations (`tools/xai_study/`). The repository also holds the reproducibility appendix that maps every reported number to its committed record and run command. A versioned snapshot will be archived on Zenodo with a minted DOI. An optional Code Ocean capsule reproduces the leaderboard and the figures from the committed result records alone. The benchmark rests on a differentiable, bit-exact emulator. It has two interchangeable versions (`jutari` in Julia, `jxtari` in JAX/Python), and both are already public. The emulator is described in the companion port paper [4]

(arXiv:2606.22447), and its source is openly available at the same repository. The companion paper supplies the bit-exactness validation and the gradient caveat that the present work imports. What is new here is the interpretability benchmark, the oracle scoring, and the semantic-gap diagnosis. We release the code under a permissive open-source licence so that a third party can score a new method against the same ground-truth oracle.

Author contributions. A.M. conceived the study, designed the ground-truth audit and the $F \wedge S \wedge M$ triad, implemented the benchmark and oracle scoring, and wrote the manuscript. S.B. contributed to the method design and the analysis and revised the manuscript. P.K. contributed to the differentiable substrate and oracle and to the gradient-faithfulness analysis, and helped develop the differentiability framing. All authors discussed the results, interpreted the semantic-gap finding, and approved the final manuscript.

Competing interests. The authors declare no competing interests.

Ethics and broader impact. This study uses no human-subjects data; the plausibility axis is a documented plausibility proxy by method tradition rather than a measurement, and we report it as such. The subject is a fully deterministic, hand-coded program with no learned policy, so no agent is trained. Interpretability methods are dual-use: a tool that audits a system can also be used to probe it. Our contribution is a calibration that makes explanation claims more trustworthy, and we release it openly to that end. The game cartridges are not redistributed.

Acknowledgements. The authors thank the maintainers of the AtariARI [6] and OCArari [7] projects, whose game-concept labels we imported as external candidates and only causally verified, and the maintainers of the `xitari` reference emulator [5] against which the substrate is validated bit-for-bit. The pilots were run locally; the full sweeps used the institutional high-performance compute cluster of Friedrich-Alexander-Universität Erlangen-Nürnberg. This work received no specific external funding.

References

- [1] Greydanus, S., Koul, A., Dodge, J. & Fern, A. *Visualizing and understanding Atari agents. Proceedings of the 35th International Conference on Machine Learning (ICML)*, Vol. 80, 1792–1801 (PMLR, 2018).
- [2] Jonas, E. & Kording, K. P. Could a neuroscientist understand a microprocessor? *PLOS Computational Biology* **13**, e1005268 (2017).
- [3] Atrey, A., Clary, K. & Jensen, D. *Exploratory not explanatory: Counterfactual analysis of saliency maps for deep reinforcement learning. International Conference on Learning Representations (ICLR)* (2020).
- [4] Maier, A., Bayer, S. & Krauss, P. A differentiable Atari VCS: A complex, fully known ground truth for explainable AI. arXiv:2606.22447; code at

- <https://github.com/akmaier/UnderstandingVCS> (2026). Companion port paper (Paper 1) of this program; submitted 2026-06-21. Code: <https://github.com/akmaier/UnderstandingVCS>. Citation key `maier2025vcs` retained for stability., 2606.22447.
- [5] DeepMind. Xitari: An Arcade Learning Environment fork. <https://github.com/google-deepmind/xitari> (2016). Accessed: 2026-06-16.
 - [6] Anand, A. *et al.* *Unsupervised state representation learning in atari*. *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 32 (2019).
 - [7] Delfosse, Q., Blüml, J., Gregori, B., Sztwiertnia, S. & Kersting, K. OCArari: Object-centric Atari 2600 reinforcement learning environments (2023). 2306.08649.
 - [8] Hewitt, J. & Liang, P. *Designing and interpreting probes with control tasks*. *Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2733–2743 (2019).
 - [9] IEEE. IEEE standard glossary of software engineering terminology (ieee std 610.12-1990). IEEE, New York (1990). Defines “software” as computer programs, procedures, and possibly associated documentation and data.
 - [10] Simonyan, K., Vedaldi, A. & Zisserman, A. *Deep inside convolutional networks: Visualising image classification models and saliency maps*. *International Conference on Learning Representations (ICLR) Workshop* (2014).
 - [11] Sundararajan, M., Taly, A. & Yan, Q. *Axiomatic attribution for deep networks*. *Proceedings of the 34th International Conference on Machine Learning (ICML)*, Vol. 70, 3319–3328 (PMLR, 2017).
 - [12] Smilkov, D., Thorat, N., Kim, B., Viégas, F. & Wattenberg, M. *Smoothgrad: Removing noise by adding noise*. *ICML Workshop on Visualization for Deep Learning* (2017).
 - [13] Springenberg, J. T., Dosovitskiy, A., Brox, T. & Riedmiller, M. *Striving for simplicity: The all convolutional net*. *International Conference on Learning Representations (ICLR) Workshop* (2015).
 - [14] Selvaraju, R. R. *et al.* Grad-CAM: Visual explanations from deep networks via gradient-based localization. *International Journal of Computer Vision* **128**, 336–359 (2020). Originally arXiv:1610.02391 (2016).
 - [15] Chattopadhyay, A., Sarkar, A., Howlader, P. & Balasubramanian, V. N. *Grad-CAM++: Generalized gradient-based visual explanations for deep convolutional networks*. *IEEE Winter Conference on Applications of Computer Vision (WACV)*, 839–847 (IEEE, 2018).

- [16] Zeiler, M. D. & Fergus, R. *Visualizing and understanding convolutional networks*. *European Conference on Computer Vision (ECCV)*, 818–833 (2014).
- [17] Fong, R. C. & Vedaldi, A. *Interpretable explanations of black boxes by meaningful perturbation*. *IEEE International Conference on Computer Vision (ICCV)*, 3429–3437 (2017).
- [18] Petsiuk, V., Das, A. & Saenko, K. *RISE: Randomized input sampling for explanation of black-box models*. *British Machine Vision Conference (BMVC)* (2018).
- [19] Ribeiro, M. T., Singh, S. & Guestrin, C. *“why should i trust you?”: Explaining the predictions of any classifier*. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, 1135–1144 (2016).
- [20] Lundberg, S. M. & Lee, S.-I. *A unified approach to interpreting model predictions*. *Advances in Neural Information Processing Systems (NeurIPS)* (2017).
- [21] Such, F. P. *et al.* *An Atari model zoo for analyzing, visualizing, and comparing deep reinforcement learning agents*. *Proceedings of the 28th International Joint Conference on Artificial Intelligence (IJCAI)*, 3260–3267 (2019).
- [22] Vig, J. *et al.* *Investigating gender bias in language models using causal mediation analysis*. *Advances in Neural Information Processing Systems (NeurIPS)* (2020).
- [23] Meng, K., Bau, D., Andonian, A. & Belinkov, Y. *Locating and editing factual associations in GPT*. *Advances in Neural Information Processing Systems (NeurIPS)* (2022).
- [24] Wang, K., Variengien, A., Conmy, A., Shlegeris, B. & Steinhardt, J. *Interpretability in the wild: A circuit for indirect object identification in GPT-2 small*. *International Conference on Learning Representations (ICLR)* (2023).
- [25] Conmy, A., Mavor-Parker, A. N., Lynch, A., Heimersheim, S. & Garriga-Alonso, A. *Towards automated circuit discovery for mechanistic interpretability*. *Advances in Neural Information Processing Systems (NeurIPS)* (2023).
- [26] Geiger, A., Lu, H., Icard, T. & Potts, C. *Causal abstractions of neural networks*. *Advances in Neural Information Processing Systems (NeurIPS)* (2021).
- [27] Geiger, A., Wu, Z., Potts, C., Icard, T. & Goodman, N. D. *Finding alignments between interpretable causal variables and distributed neural representations*. *Proceedings of the Third Conference on Causal Learning and Reasoning (CLeaR)* (2024).
- [28] Cunningham, H., Ewart, A., Riggs, L., Huben, R. & Sharkey, L. *Sparse autoencoders find highly interpretable features in language models* (2023). [2309.08600](#).

- [29] Alain, G. & Bengio, Y. *Understanding intermediate layers using linear classifier probes. International Conference on Learning Representations (ICLR) Workshop* (2017).
- [30] Belrose, N. *et al.* Eliciting latent predictions from transformers with the tuned lens (2023). [2303.08112](#).
- [31] Nanda, N. Attribution patching: Activation patching at industrial scale. <https://www.neelnanda.io/mechanistic-interpretability/attribution-patching> (2023).
- [32] Karvonen, A. *et al.* *SAEBench: A comprehensive benchmark for sparse autoencoders in language model interpretability. Proceedings of the 42nd International Conference on Machine Learning (ICML)*, Vol. 267 of *PMLR* (2025). Eight metrics over >200 open-sourced SAEs; finds gains on proxy metrics do not reliably translate to practical interpretability., [2503.09532](#).
- [33] Jacovi, A. & Goldberg, Y. *Towards faithfully interpretable NLP systems: How should we define and evaluate faithfulness? Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, 4198–4205 (Association for Computational Linguistics, 2020).
- [34] Adebayo, J. *et al.* *Sanity checks for saliency maps. Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 31, 9505–9515 (2018).
- [35] Paper 2 authors. Faithful-method demonstration: causal near-ceiling versus popular near-chance on the differentiable Atari VCS. Committed analysis record, `tools/xai_study/compare/out/faithful_demo.md`; pure read over the committed cross-tradition leaderboard (`compare/out/leaderboard.json`) and the per-game §R records it aggregates. Generated by `compare/faithful_demo.py`. (2026). Internal data artifact. Position-regime pair `activation_patching = 1.000` vs `vanilla_saliency = 0.000` on all six core games (per-method gap 1.000). Accessed 2026-06-23.
- [36] Paper 2 authors. Cross-tradition faithfulness–plausibility leaderboard for the differentiable Atari VCS interpretability audit (phases a–c). Committed analysis record, `tools/xai_study/compare/out/leaderboard.json`; aggregates the per-game §R records under `tools/xai_study/phaseA_kording/out/` (2026). Internal data artifact; 257 per-game records aggregated into 31 method rows including the oracle positive control. Generated by `compare/leaderboard.py`. Accessed 2026-06-23.
- [37] Yang, M. & Kim, B. BAM: Benchmarking attribution methods with relative feature importance (2019). Introduces the BAM benchmark: models trained with known relative feature importance plus paired-model/paired-input metrics for attribution methods. (Key `yang2019bim` retained for citation stability; the work is BAM, not BIM.), [1907.09701](#).

- [38] Li, X. *et al.* $\mathcal{M}^4$: A unified XAI benchmark for faithfulness evaluation of feature attribution methods across metrics, modalities and models. *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, Vol. 36 (2023).
- [39] Lindner, D. *et al.* Tracr: Compiled transformers as a laboratory for interpretability. *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 36 (2023). [2301.05062](#).
- [40] Gupta, R., Arcuschin, I., Kwa, T. & Garriga-Alonso, A. *InterpBench: Semi-synthetic transformers for evaluating mechanistic interpretability techniques. Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, Vol. 37 (2024). [2407.14494](#).
- [41] Mueller, A. *et al.* MIB: A mechanistic interpretability benchmark. *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, Vol. 267 of *PMLR* (2025). Two tracks (circuit localization; causal-variable localization) over four tasks and five neural language models., [2504.13151](#).
- [42] Bellemare, M. G., Naddaf, Y., Veness, J. & Bowling, M. The Arcade Learning Environment: An evaluation platform for general agents. *Journal of Artificial Intelligence Research* **47**, 253–279 (2013).
- [43] Machado, M. C. *et al.* Revisiting the Arcade Learning Environment: Evaluation protocols and open problems for general agents. *Journal of Artificial Intelligence Research* **61**, 523–562 (2018).
- [44] Dalton, S., Frosio, I. & Garland, M. *Accelerating reinforcement learning through GPU Atari emulation. Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 33 (2020). ArXiv:1907.08467; code at <https://github.com/NVlabs/cule>.
- [45] Delfosse, Q. *et al.* JAXAtari: A high-performance framework for reasoning agents. <https://github.com/k4ntz/JAXAtari> (2026). Software; accessed 2026-06-19.
- [46] Barbiero, P. *et al.* Neural interpretable reasoning. *arXiv preprint arXiv:2502.11639* (2025).
- [47] Marr, D. *Vision: A Computational Investigation into the Human Representation and Processing of Visual Information* (W. H. Freeman and Company, 1982).
- [48] Lazebnik, Y. Can a biologist fix a radio? — or, what i learned while studying apoptosis. *Cancer Cell* **2**, 179–182 (2002).
- [49] Bengio, Y., Léonard, N. & Courville, A. Estimating or propagating gradients through stochastic neurons for conditional computation (2013). [1308.3432](#).

- [50] Jang, E., Gu, S. & Poole, B. *Categorical reparameterization with Gumbel-softmax*. *International Conference on Learning Representations (ICLR)* (2017).
- [51] Jaderberg, M., Simonyan, K., Zisserman, A. & Kavukcuoglu, K. *Spatial transformer networks*. *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 28 (2015).
- [52] Abnar, S. & Zuidema, W. *Quantifying attention flow in transformers*. *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*, 4190–4197 (2020).
- [53] Chan, L. *et al.* Causal scrubbing: A method for rigorously testing interpretability hypotheses. *AI Alignment Forum* (2022).
- [54] Shrikumar, A., Greenside, P. & Kundaje, A. *Learning important features through propagating activation differences*. *Proceedings of the 34th International Conference on Machine Learning (ICML)*, 3145–3153 (2017).
- [55] Erion, G., Janizek, J. D., Sturmfels, P., Lundberg, S. M. & Lee, S.-I. Improving performance of deep learning models with axiomatic attribution priors and expected gradients. *Nature Machine Intelligence* **3**, 620–631 (2021).
- [56] Štrumbelj, E. & Kononenko, I. Explaining prediction models and individual predictions with feature contributions. *Knowledge and Information Systems* **41**, 647–665 (2014).
- [57] Fong, R., Patrick, M. & Vedaldi, A. *Understanding deep networks via extremal perturbations and smooth masks*. *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, 2950–2958 (2019).
- [58] Olson, M. L., Khanna, R., Neal, L., Li, F. & Wong, W.-K. Counterfactual state explanations for reinforcement learning agents via generative deep learning. *Artificial Intelligence* **295**, 103455 (2021).
- [59] Bastani, O., Pu, Y. & Solar-Lezama, A. *Verifiable reinforcement learning via policy extraction*. *Advances in Neural Information Processing Systems (NeurIPS)*, Vol. 31, 2494–2504 (2018).
- [60] Balduzzi, D. *et al.* *The shattered gradients problem: If resnets are the answer, then what is the question?* *Proceedings of the 34th International Conference on Machine Learning (ICML)*, Vol. 70 of *PMLR*, 342–350 (2017).
- [61] Morcos, A. S., Barrett, D. G. T., Botvinick, M. & Rabinowitz, N. C. *On the importance of single directions for generalization*. *International Conference on Learning Representations (ICLR)* (2018). [1803.06959](#).

- [62] Leavitt, M. L. & Morcos, A. S. *Towards falsifiable interpretability research. NeurIPS 2020 Workshop: ML Retrospectives, Surveys & Meta-Analyses (ML-RSA)* (2020). [2010.12016](#).
- [63] Elhage, N. *et al.* Toy models of superposition (2022). Transformer Circuits Thread, https://transformer-circuits.pub/2022/toy_model/index.html, [2209.10652](#).
- [64] Kindermans, P.-J. *et al.* in *The (un)reliability of saliency methods Explainable AI: Interpreting, Explaining and Visualizing Deep Learning*, Vol. 11700 of *Lecture Notes in Computer Science* 267–280 (Springer, 2019).
- [65] Sturmfels, P., Lundberg, S. & Lee, S.-I. Visualizing the impact of feature attribution baselines. *Distill* **5** (2020). <https://distill.pub/2020/attribution-baselines>.
- [66] Jain, S. & Wallace, B. C. *Attention is not explanation. Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, 3543–3556 (Association for Computational Linguistics, 2019).

A taxonomy of interpretability failure modes on the VCS, by underlying mechanism




Fig. 6 A taxonomy of interpretability failure on the VCS, organised by mechanism. Each leaf is a method exhibiting the branch’s failure mode, annotated with its *family-mean* faithfulness $F \pm 95\%$ bootstrap-over-games CI against the intervention oracle (Section 3.2; 0 = chance, 1 = oracle ceiling), read from `leaderboard.ci.csv`; the per-game exemplars quoted below carry their own “on (game)” label. The lower branches confirm the expected failures: the gradient family is provably zero on the discrete position/index regime, and the architecture-specific methods (Grad-CAM, attention, VIPER) do not apply because the VCS lacks the neural substrate they require, a property of the subject, not a method failure. The higher-order branches are the *semantic gap*: present-but-unused concepts (linear probing [8] flags a present-but-unused cell the oracle says is inert), named-but-inert factorisations (the SAE [28] matches every variable at fraction 1.0 yet scores $F = 0.041$ on Pong), and graph-correct-but- incomplete circuits (ACDC [25] at $F = M = 1.0$, $S = 0.44$ on Breakout). The neural-network parallels shown alongside each branch are documented analogues [60–66], evidence transfer rather than proof. Every annotated value reads from a committed record.

Every failure mode measured on the VCS has a documented neural-network analogue

The map reads as a necessary-condition screen: on the VCS each failure is proven against the exact §3.2 intervention oracle (full observability, no learning); on neural networks the same failure is a documented analogue, not proven here. A method that fails the screen should not be trusted for stronger claims without further evidence (necessary, not sufficient).

Failure mode measured on the VCS	VCS measured evidence proven in this benchmark	Neural-network analogue documented, not proven here	Status here → there
1 Gradient vanishes on discrete / index outputs	0.000 position-regime F single record Vanilla saliency & IG = 0.000 on the position-regime (provably zero at a strobe-timed argmax); whole gradient bucket 0.068.	Shattered / saturated gradients; non-differentiable argmax & index readouts.	proven here analogue there
2 Saliency is model-invariant (sanity-check fails)	0.000 sanity-check margin single record Program-randomization: saliency cosine = 1.000 while 95% of RAM changed -> sanity check FAILS.	Saliency unchanged under model- & label-randomization; behaves like an edge detector.	proven here analogue there
3 Correlation ≠ causation (Granger & tuning trap)	0.136 Granger F 95% CI [0.02, 0.29] Granger-edge faithfulness = 0.136 (false edges vs the true data-flow); game-variable tuning = 0.116.	Spurious Granger edges; 'present ≠ used' tuning & attention read as explanation.	proven here analogue there
4 Single-unit ablation has an interaction blind-spot	0.414 connectome F 95% CI [0.08, 0.78] Single-unit lesion is faithful (0.987) yet misses the wiring: connectome recovered at only 0.414.	Single-unit ablation misleading: computation is distributed across many units.	proven here analogue there
5 SAE / probe: present-not-used + polysemanticity	0.162 selectivity F 95% CI [0.08, 0.24] SAEs match a known variable yet score 0.195 on causal use: probes decode at 0.894 acc but 25 cells are not causal.	Control-task selectivity gap: superposition / polysemantic neurons.	proven here analogue there
6 Baseline-sensitivity of Integrated / Expected Gradients	0.034 EG F (all-regime) 95% CI [0.00, 0.10] IG magnitude is baseline-sensitive (up to 0.999 across games); a content-baseline kills IG; EG lands at 0.034.	Attribution depends on the (arbitrary) baseline choice; no canonical reference input.	proven here analogue there

Headline (position regime): causal/intervention 0.412 vs gradient/correlational 0.068 (95% CI 0.390 / 0.070)
 VCS evidence: proven in this benchmark (vs the exact §3.2 intervention oracle); 0.344 faithfulness gap; activation patching = 1.000 (oracle ceiling) where vanilla saliency = 0.000 (naive-gradient floor).
 NN analogue: documented in the literature, not proven by this benchmark

Fig. 7 The VCS↔neural-network failure-mode map. Each row pairs a failure mode we measured on the VCS against the oracle (§3.2) — left, blue — with the neural-network behavior it is a documented analogue of — right, vermillion. On the VCS the failure is provable because we hold the exact ground truth; the neural-network column is a documented analogue, not proven by this benchmark, and reads the attribution and probing literature ([11, 34, 60, 65] on gradients; [8, 61, 62, 66] on use vs. presence; [6, 63, 64] on probing, superposition, and reliability). The takeaway is one line: every VCS failure has a named counterpart in how interpretability methods are reported to fail on networks, so a property provable here is one to suspect there. Per-row VCS scores, their committed records, and the full map ([fig5.representativeness_map_full.pdf](#)) are in the Supplementary Information.

Supplementary Information

This Supplementary Information records what the main paper summarizes. We give the per-method protocols (Sec. S1) and the leaderboard schema with its glossary (Sec. S2). We give the coverage and applicability of every method across games and output regimes (Sec. S3). We map each headline claim to its evidence (Sec. S4). We trace the provenance of every headline number down to the record file and command that produced it (Sec. S5). We state plainly what the paper does not claim (Sec. S6). We give the plausibility-proxy rubric and its sensitivity analysis (Sec. S7). The main text shows simplified forms of two diagnostic figures; we reproduce both here at full resolution (Figs. S1 and S2). The correctness triad $F \wedge S \wedge M$ used throughout is the one defined in the main Methods (Section 3.4); we reference it here rather than restating it. The benchmark scores 30 interpretability methods together with the intervention oracle (Section 3.2) as a positive control, for 31 leaderboard rows in total.

S1 Per-method protocols

This section records five things for each method under test. It records the exact objects the method consumes and produces, the reference it is scored against, its sampling or optimization budget, the metric that turns its output into a faithfulness score, and the committed record file that holds the per-game results. All methods share one evaluation contract. We state that contract once and then specialize it per method. Every method receives the same bit-exact VCS trajectory on one of the six core games (`breakout`, `ms_pacman`, `pong`, `qbert`, `seaquest`, `space_invaders`). The candidate-cause set is the machine’s own state. This set is the 128 RAM bytes together with the CPU registers and the TIA and RIOT registers exposed by the substrate. The reference is the intervention oracle of Section 3.2, evaluated by exact re-runs of the emulator under $\text{do}(u)$. A method that cannot be applied to a given output regime is recorded as not-applicable rather than scored as zero. An empty result is therefore never read as a wrong one. Faithfulness is the F axis of the main-text triad (Section 3.4); we reference that definition and do not restate it. Records live under `tools/xai_study/<phase>/out/`; the aggregate they feed is `tools/xai_study/compare/out/leaderboard.json`.

The neuroscience battery (Phase A) treats the running machine as a brain. It asks how far each classical analysis recovers the known mechanism. Each method’s baseline is the exact data-flow graph or causal role that the substrate makes available. Its budget is the six-game core unless noted, and its scoring metric is named in Table S1.

The attribution battery (Phase B) runs the everyday XAI toolbox over the same trajectories. These methods produce a real-valued saliency or attribution over the candidate-cause set. We score that output by Pearson correlation against the oracle attribution. The gradient-based members also inherit the gradient-vanishing failure of the discrete position regime (Prop. 1). The budgets that matter are the Monte-Carlo or perturbation counts of the sampling estimators, recorded per method in Table S2. The seed-dispersion of the stochastic members is carried in `leaderboard.ci.csv` (`seed_sd`, `seed_sd_source`) and reproduced in Sec. S5.

Table S1 Phase A protocols (neuroscience battery). One row per method. The columns give the input it reads, the reference it is scored against, the sampling or intervention budget, the scoring metric, and the committed record. All eight run on the six core games; faithfulness is the F axis of Section 3.4. The whole-state recording (A8) is included as a descriptive upper bound on coverage. Its faithfulness F is trivial, so we read it for its minimality M .

Method	Input / candidate causes	Reference	Scoring metric	Record
A1 connec-tomics	RAM/register activity over trajectory	data-flow graph	data-flow graph F_1 vs. oracle	A1
A2 lesions	single-cell knock-outs via $\text{do}(u)$	true causal role	lesion-importance ρ vs. role	A2
A3 tuning curves	per-cell response vs. variable	variable coupling	1– spurious-tuning rate	A3
A4 spike-word / pairwise	pairwise cell correlations	true coupling	corr. ρ vs. coupling	A4
A5 field potentials	pooled-activity spectra	true causal share	1– clock-explained var.	A5
A6 Granger	directed lag structure	true data-flow	1– false-edge rate	A6
A7 dim. reduction	NMF/PCA over activity	known variables	NMF matched-component frac.	A7
A8 whole-state	full state record	oracle decomp.	minimality M vs. oracle	A8

Records: `tools/xai_study/phaseA_kording/out/A{1..8}_*`; all eight run on the six core games.

The mechanistic battery (Phase C) runs the methods built to recover circuits. Each is scored against the true data-flow or the exact recovered-minus-exact deviation. The causal members of this phase are the ones that reach the oracle ceiling. We list their exact scoring metric so that the ceiling is not read as a tautology. Each member is scored against the substrate’s known routine, and not against another method. Two members (NMF/PCA dictionaries and the sparse autoencoder) are dimensionality-reduction methods carried here for comparison. The sparse autoencoder has two aggregations that differ, which we reconcile in Sec. S5.

The intervention oracle is the positive control, not a method under test. It is the exact $\text{do}(u)$ ground truth of Section 3.2, evaluated by re-running the bit-exact machine. It attains $F = S = M = 1.0$ by construction. Its record triple is the intervention oracle, the content-path gradient oracle, and their cross-check (`tools/xai_study/ground_truth/out/`). It anchors the ceiling against which every other row is read, which is why it occupies the thirty-first leaderboard row.

S2 Benchmark schema and glossary

The leaderboard is built from a flat record schema. Stating it removes the ambiguity about what a single number on the benchmark counts. Each method is run per game and writes a per-game record under its phase’s `out/` tree. The leaderboard aggregates those records into one row per method by averaging over games. The schema columns

Table S2 Phase B protocols (attribution battery). Each method emits a saliency over the candidate-cause set. We score it by Pearson correlation against the oracle attribution (Section 3.2). The reference and the six-game budget are shared. The budget column records the sampling or perturbation count that drives each estimator. The seed column flags which methods carry an independent-seed dispersion in `leaderboard.ci.csv`.

Method	Tradition	Budget / hyperparameters	Seed sd
<code>vanilla_saliency</code>	gradient	single backward pass	—
<code>smoothgrad</code>	gradient	noise-averaged gradient; σ -robustness committed	0.0
<code>guided_backprop</code>	gradient	single guided pass	—
<code>gradxinput</code>	gradient	DeepLIFT reference-difference	—
<code>_deeplift</code>			
<code>integrated</code>	gradient	path integral, baseline = zero	—
<code>_gradients</code>		state	
<code>expected</code>	gradient	expectation over baselines	yes
<code>_gradients</code>		(refits)	
<code>kernelshap</code>	gradient	Monte-Carlo coalition sampling ($\propto 1/\sqrt{N}$)	0.017
<code>lime</code>	gradient	local linear fit, perturbation refits	0.029
<code>rise</code>	gradient	random-mask Monte-Carlo ($\propto 1/\sqrt{N}$)	0.051
<code>occlusion</code>	intervention	exhaustive single-cell masking	—
<code>extremal</code>	intervention	optimized mask, area-constrained	—
<code>_perturbation</code>			
<code>on_distribution</code>	intervention	on-manifold $\text{do}(u)$ resampling	—
<code>_counterfactual</code>			

Records: `tools/xai_study/phaseB.attribution/out/`; all twelve run on the six core games and are scored by Pearson correlation against the oracle (Section 3.2). Seed sd is the independent-seed dispersion from `leaderboard.ci.csv`.

of a leaderboard row are given in Table S4. The aggregation rule is a mean over the n_{games} games. When a method emits several output records on one game, those records are first reduced within that game. The faithfulness, sufficiency, and minimality columns report the F, S, and M axes defined in Section 3.4. We do not re-derive those definitions here.

We fix the vocabulary that the schema implies, because the review asked for a single unambiguous method count. A *method* is one interpretability technique, e.g. `vanilla_saliency`. A *method row* is its single aggregated line on the leaderboard. A *per-game record* is one method scored on one game, written under a phase `out/` tree. An *output record* is a per-game record for one output regime; a method may emit several of these on one game. A *family mean* is the average over the members of a tradition family, e.g. causal/intervention. The *oracle positive control* is the exact $\text{do}(u)$ ground truth of Section 3.2, scored as a method to anchor the ceiling. A *not-applicable method* is one that does not apply to a given output regime; it is recorded as N/A rather than as a zero. With this vocabulary the count is unambiguous. The benchmark scores 30 interpretability methods together with the intervention oracle as a positive control, for 31 leaderboard rows in total (`leaderboard.json.n.methods` = 31), aggregating 257 per-game records (`n_per_game_records.aggregated` = 257): 8

Table S3 Phase C protocols (mechanistic battery). Each method is scored against the substrate’s known circuit or the exact recovered-minus-exact deviation (Section 3.2), and not against another method. This is what keeps the oracle ceiling non-circular. The probing and dictionary-learning members are included for comparison with the causal members.

Method	Tradition	Scoring metric
activation_patching	causal	$\max \text{recovered} - \text{exact} $ vs. oracle
causal_scrubbing	causal	scrubbing-preserved performance
interchange	causal	aligned interchange accuracy
_interventions_das		
logit_tuned_lens	causal	logit-lens fidelity to true intermediate
path_patching	causal	path-circuit F_1 vs. true routine
ACDC	causal	discovered-circuit best F_1 vs. data-flow
attribution_patching	gradient	corr. of approximation vs. exact patch
linear_probing	probing	selectivity vs. control task
_control_tasks		
nmf_pca_dictionaries	dim. red.	matched-component fraction vs. vars
sparse_autoencoder	dim. red.	feature-variable matched fraction

Records: `tools/xai_study/phaseC_mechanistic/out/`; all ten run on the six core games.

methods in Phase A, 12 in Phase B, 10 in Phase C, and the one oracle row. This count is the canonical one used in the abstract, the figures, and throughout this Supplement.

S3 Coverage and applicability

The benchmark scores methods on a six-game core drawn from a larger bit-exact substrate. The coverage table makes the scope of that choice explicit. Paper 1 validates the differentiable VCS ports bit-for-bit against the `xitari` reference on all 64 Arcade Learning Environment games, within a 30-frame RAM and 60-frame screen window. The interpretability study reported here runs on the six games listed in Table S5. We pick these six because all three tiers of ground truth (Section 3.3) are available for them and the oracle is exact. The remaining 58 games are bit-exact substrate but are not scored here. Some fall outside the validated horizon for the analyses we run. For others, the tiered labels they would need are not yet produced. They are candidates for the larger sweep, and are not excluded for any methodological reason. The horizon column records the validated window. The tier column records which of the three ground-truth tiers (T1 exact intervention, T2 data-flow, T3 externally-supplied semantic labels) is available.

Whether a method applies at all depends on what the method needs from its target. The applicability matrix records that for every method family. The Phase A neuroscience analyses read the raw machine state and need neither a learned policy nor gradients. The Phase B attribution methods split into two groups. The gradient members need a differentiable forward pass. The intervention members need only the ability to set state. The Phase C mechanistic methods need access to internal activations. The causal members among them also need the ability to intervene. Table S6

Table S4 Leaderboard row schema. The columns of one row in `compare/out/leaderboard.json (rows[])` and its uncertainty companion `leaderboard.ci.csv`. The metric column names the per-method scoring function from Sec. S1; F, S, and M are the triad axes of Section 3.4. Aggregation is a mean over n_{games} , with bootstrap confidence intervals over games (and, where present, over records).

Column	Meaning
<code>phase</code>	<code>phaseA_kording</code> , <code>phaseB_attribution</code> , <code>phaseC_mechanistic</code> , or <code>ground_truth</code>
<code>method</code>	method identifier (the row key)
<code>tradition</code>	intervention, causal, gradient, correlational, <code>dim_reduction</code> , probing, descriptive, oracle
<code>n_games</code>	number of games the method was scored on (6 for the core set; 1 for the oracle)
<code>n_records</code>	number of per-game output records aggregated into the row
<code>faithfulness (F)</code>	faithfulness vs. the intervention oracle (Section 3.2), $[0, 1]$, higher better
<code>faithfulness_ci95</code>	95% bootstrap CI half-width over games
<code>faithfulness_position_regime</code>	F restricted to the discrete sprite-position/index regime
<code>faithfulness_content_regime</code>	F restricted to the continuous content regime
<code>triad_S / triad_M</code>	sufficiency / minimality axes (Section 3.4), where applicable
<code>plausibility_proxy</code>	documented method-tradition proxy, $[0, 1]$ — not a human-subjects measurement
<code>is_positive_control</code>	1 for the oracle row, 0 otherwise
<code>metric_names</code>	the per-method scoring function (Sec. S1)
<code>commits</code>	the commit(s) at which the per-game records were produced

Table S5 Core-game coverage. The six core games carry exact tiered ground truth and the full method battery. Faithfulness is scored within the Paper 1 bit-exact window. The last two columns give the methods run and the per-game output records contributed. The wider 64-game substrate is bit-exact (Paper 1) but is reserved for the larger sweep. T3 labels are external (Section 3.3); we record this so the benchmark does not appear to certify them.

Game	Paper-1 status	Horizon	Tiers	Methods run	Records
<code>breakout</code>	bit-exact	30/60 fr	T1,T2,T3	30 + oracle	per-game
<code>ms_pacman</code>	bit-exact	30/60 fr	T1,T2,T3	30 + oracle	per-game
<code>pong</code>	bit-exact	30/60 fr	T1,T2,T3	30 + oracle	per-game
<code>qbert</code>	bit-exact	30/60 fr	T1,T2,T3	30 + oracle	per-game
<code>seaquest</code>	bit-exact	30/60 fr	T1,T2,T3	30 + oracle	per-game
<code>space_invaders</code>	bit-exact	30/60 fr	T1,T2,T3	30 + oracle	per-game
58 further games	bit-exact (P1)	30/60 fr	T1,T2 partial	reserved	—

records, per family, whether the method operates on the raw VCS, whether it requires a learned policy, gradients, internal layers or attention, or external labels, whether it uses oracle-like interventions, and whether it receives a supplied hypothesis. The table shows the same pattern the results confirm. The families that score high on faithfulness are the ones that can intervene. The families that collapse on the position regime are the gradient and correlational ones.

Table S6 Per-family applicability. What each method family requires of its target. “Intervenes” marks oracle-like $\text{do}(u)$ access; “hypothesis” marks methods that receive a supplied circuit hypothesis to test. The raw-VCS column distinguishes analyses that read machine state directly from those that need a differentiable or layered model. A blank means the requirement does not apply.

Family	raw VCS	policy	gradients	layers/attn	intervenes	hypothesis
A neuroscience	yes	no	no	no	A2 only	no
B gradient	yes	no	yes	no	no	no
B intervention	yes	no	no	no	yes	no
C causal	yes	no	no	yes	yes	yes
C gradient (attr. patch)	yes	no	yes	yes	approx.	yes
C probing	yes	no	no	yes	no	no
C / A dim. reduction	yes	no	no	yes	no	no
oracle (control)	yes	no	no	yes	exact	—

S4 Claims and evidence

Every headline claim in the paper rests on a specific artifact. The table in this section makes that dependency checkable. For each claim we list the figure or table that supports it, the games and methods it uses, the metric, the artifact it traces to, and the limitation that bounds it. Faithfulness, sufficiency, and minimality are the three axes of Section 3.4. The numbers themselves are sourced in Sec. S5. The table below names which evidence supports which claim; it does not repeat the values.

S5 Number provenance

Every headline number traces to a committed record, the script that produced it, and the command that re-derives it. This section is the in-paper copy of that traceability table. It mirrors Section 6.5 of the reproducibility document (`xai_2.interpretability/REPRODUCIBILITY.md`). The values are the measured, committed ones at the current revision. The confidence intervals are read from `tools/xai_study/compare/out/leaderboard_ci.csv` rather than recomputed here. Each record also carries `seed`, `where`, `commit`, and `timestamp`, so the provenance of every value is self-documenting.

One number needs a reconciliation, and we record it rather than hide it. The sparse autoencoder is the only method whose two aggregations disagree. The over-games mean is $F = 0.1404$ with a 95% bootstrap CI of $[0.0, 0.3139]$. The over-records mean is 0.1947 with CI $[0.0058, 0.4804]$. Both are read from `leaderboard_ci.csv`. The figures and the main-text leaderboard plot the over-records value 0.1947 , which lies inside the over-games interval. We treat the over-games mean as the canonical row value for the aggregation rule of Sec. S2, and report the over-records figure alongside it, so that text, table, and figure agree on which aggregation each cites. No headline claim depends on which of the two is read, since the sparse autoencoder sits near the floor under both.

Table S7 Claims and their evidence. Each headline claim maps to the figure or table that supports it, the games and methods involved, the metric, the committed artifact, and the limitation that bounds the claim. The artifact column points to the record that Sec. S5 traces to a value. The limitation column states the scope, so the claim is not read wider than the evidence.

Claim	Evidence	Methods / games	Metric	Limitation
Causal/intervention methods beat gradient/correlational on faithfulness across all regimes	Fig. 4; leaderboard	25 methods, 6 games	F mean (gap 0.387)	not robust: gap collapses to 0.123 $[-0.008, 0.256]$ excl. oracle-like and 0.045 $[-0.087, 0.153]$ content-only (CIs cross zero); only the position-regime gap 0.3435 is robust
The gap holds with the same sign on the discrete position regime	Fig. 4; faithful_demo	13 rows, 6 games	F position (gap 0.344)	position regime only
A faithful causal method reaches the oracle ceiling where the default tool collapses	faithful_demo ; Fig. 4	act. patching vs. vanilla saliency, 6 games	F 1.0 vs. 0.0 (position regime; all-regimes saliency F 0.267, plotted in Fig. 4)	one decisive pair, position regime only
The naive gradient of a discrete index output is provably zero	Prop. 1 (Section 3.2); oracle_grad	all gradient methods	grad@index = 0	index path only
High plausibility coexists with low faithfulness (the danger zone)	Fig. 4; faithful_demo	vanilla saliency	proxy 0.9, F 0.0 (position regime; all-regimes F 0.267, plotted in Fig. 4)	proxy, not measured plausibility
Faithful attribution recovers the wiring but not the computation's semantics	Fig. 6 (Section S3); discussion	Phase B vs. C	F vs. S, M	necessary-condition stress test
The failure modes have documented NN analogues	Fig. 7 (Fig. S1)	taxonomy	—	analogues, not proof

S6 What this paper does not claim

The benchmark is sharp because its scope is narrow. Stating the boundaries keeps the result from being read wider than the evidence allows. We do not claim to recover a program's semantics from scratch. The study measures whether an explanation names the true causes and regenerates the behavior on a machine whose answer is already known. It does not show that any method discovers the meaning of a routine without

that ground truth in hand. A faithful method recovers the causal structure but not the meaning of the variables, and this remaining gap is the point rather than a result to be explained away.

We do not evaluate learned agents. The subject throughout is the bit-exact VCS itself. The candidate causes are its registers, RAM cells, and program variables, not the weights of a trained policy. The neural-network parallels drawn in the discussion and in Fig. S1 are documented analogues from the literature, not claims proven on this benchmark. We label them as such, so that a reader does not take a mapped failure mode for a demonstrated one.

We do not prove that neural-network interpretability is impossible. The VCS gap is a necessary-condition stress test. A method that fails to recover the causal structure of a fully known machine cannot be trusted to recover it on an unknown one. Yet, a method that succeeds here is not thereby certified on networks. The “floor” framing is an evidence-transfer argument, not a proven lower bound on neural-network interpretability.

We do not measure human plausibility. The plausibility axis is a documented method-tradition proxy, computed from a fixed rubric rather than collected from raters, and we report it as a proxy throughout. Its job is to expose the danger zone where a convincing explanation is unfaithful, not to stand in for a human-subjects study. We also do not redistribute the ROMs. The substrate is validated against an external reference, and the records are reproducible from committed artifacts without the ROMs. The ROM provenance is given as hashes, so a third party supplies their own legally obtained copies.

S7 Plausibility-proxy rubric and its sensitivity

The plausibility axis of the leaderboard is a documented, tradition-level proxy, not a human-subjects measurement. This section records how each value was assigned. It also shows that the headline contrast does not depend on the particular numbers chosen. The Methods (Section 3.4) name two reporting axes. Faithfulness is measured on the machine against the intervention oracle of Section 3.2. The plausibility proxy stands in for how convincing a tradition’s explanations are usually taken to be. We assign the proxy once per method tradition, before any faithfulness score on our machine is computed, so it cannot be tuned to the result. Its only job is to expose the danger zone of explanations that look convincing yet are unfaithful. For that job it must be coarse, fixed in advance, and sourced rather than precise.

The rubric assigns one proxy per tradition on a fixed five-step scale. It scores two documented properties of each tradition’s typical output and takes their average. The first property is whether the explanation produces a dense, picture-like artifact that a reader sees as an account of the computation. Heatmaps and saliency-style maps over the input are read as convincing by construction. This is the perceptual-immediacy axis. The second property is whether the field has historically treated the tradition’s output as a self-evident explanation rather than a hypothesis to be tested. This is the adoption-as-explanation axis. A tradition that scores high on both, such as gradient saliency, receives a high proxy. A tradition whose output is an austere intervention

Every failure mode measured on the VCS has a documented neural-network analogue

The map reads as a necessary-condition screen: on the VCS each failure is proven against the exact §3.2 intervention oracle (full observability, no learning); on neural networks the same failure is a documented analogue, not proven here. A method that fails the screen should not be trusted for

Failure mode (stronger if measured on the VCS)	VCS measured evidence (necessary, not proven in this benchmark)	Neural-network analogue (documented, not proven here)	Status (here → there)
1 Gradient vanishes on discrete / index outputs	0.000 position-regime F single record Vanilla saliency & IG faithfulness = 0.000 on the position regime; the whole gradient/correlational bucket = 0.068 (n=9). Provably 0 at a strobe-timed argmax readout (§3.2 oracle).	Shattered / saturated gradients; non-differentiable argmax & index readouts (analogue refs G1, G2).	proven here ↓ analogue there
2 Saliency is model-invariant (sanity-check fails)	0.000 sanity-check margin single record Program-randomization test: saliency cosine = 1.000 to the real map while 95% of RAM changed → sanity check FAILS. Constant Jacobian on the content path, no ReLU.	Saliency unchanged under model- & label-randomization; behaves like an edge detector (analogue ref S1).	proven here ↓ analogue there
3 Correlation ≠ causation (Granger & tuning trap)	0.136 Granger F 95% CI [0.02, 0.29] Granger-edge faithfulness = 0.136 (false edges vs the true data-flow); game-variable tuning = 0.136 (spurious tuning) – both vs the exact §3.2 oracle.	Spurious Granger edges; 'present ≠ used' tuning & attention-as-explanation (analogue refs C1, C2).	proven here ↓ analogue there
4 Single-unit ablation has an interaction blind-spot	0.414 connectome F 95% CI [0.08, 0.78] Single-unit lesion importance is faithful (0.987) yet MISSES the wiring: the connectome graph is recovered at only 0.414; SAE-feature ablations move in 0 of the audited cells.	Single-unit ablation misleading: computation distributed across units (analogue refs D1, D2).	proven here ↓ analogue there
5 SAE / probe: present-not-used + polysemanticity	0.162 selectivity F 95% CI [0.08, 0.24] SAEs match a known variable (100% of candidates in 4/6 games, min 83%) yet score 0.195 on causal use; probes decode at 0.894 acc but 25 cells are decodable-not-causal → selectivity F = 0.162.	Control-task selectivity gap: superposition / polysemantic neurons (analogue refs P1, P2).	proven here ↓ analogue there
6 Baseline-sensitivity of Integrated / Expected Gradients	0.034 EG F (all-regime) 95% CI [0.00, 0.10] IG magnitude is baseline-sensitive (corr. sensitivity up to 0.999 across games); a baseline equal to the content byte kills IG entirely; Expected Gradients lands at 0.034.	Attribution depends on the (arbitrary) baseline choice; no canonical reference input (analogue refs B1, B2).	proven here ↓ analogue there

Headline (position regime): causal/intervention 0.412 vs gradient/correlational 0.068 (95% CI 0.390 / 0.070)
 VCS evidence: proven in this benchmark (vs the exact §3.2 intervention oracle) 0.344 faithfulness gap; activation patching = 1.000 (oracle ceiling) where vanilla saliency = 0.000 (naive-gradient floor).
 NN analogue: documented in the literature, not proven by this benchmark

Analogue references (documented in the literature, not measured here):

G1 Balduzzi et al. 2017; G2 Sundararajan et al. 2017 (IG); S1 Adebayo et al. 2018; C1 Anand et al. 2019; C2 Jain & Wallace 2019.

D1 Morcos et al. 2018; D2 Leavitt & Morcos 2020; P1 Hewitt & Liang 2019; P2 Elhage et al. 2022; B1 Kindermans et al. 2019; B2 Sturmfels et al. 2020.

All VCS numbers read from committed records (leaderboard.json, leaderboard_ci.csv, faithful_demo.json, guided_backprop / ig_baseline_sweep / sae / linear_probing core summaries); 95% CIs are bootstrap-over-games (P2-R-L).

Fig. S1 Full VCS-to-NN representativeness map. The full-density version of the main-text map (Fig. 7). Each VCS failure mode is placed against its documented neural-network analogue from the literature. The neural-network column records documented analogues, not parallels proven by this benchmark. The oracle is the intervention oracle of Section 3.2, and the underlying per-method values trace to `compare/out/leaderboard.json` via Sec. S5. Note that the simplified main-text version condenses this full map.

table or a circuit edge list, rarely mistaken for a finished story, receives a low one. The scale is deliberately blunt, {0.3, 0.5, 0.55, 0.75, 0.8, 0.85, 0.9}. A proxy meant to expose a danger zone gains nothing from false precision, and a finer scale would invite over-reading. Every value in Table S9 traces to the `plausibility_proxy` field of the per-method rows in `compare/out/leaderboard.json`. The table only groups those rows by tradition and records the rationale.

The paper draws a contrast between this proxy and measured faithfulness. On the committed records the two move in opposite directions. Across the thirty methods under test the Spearman rank correlation between faithfulness and the proxy is -0.64 , a strong negative association. The traditions the field finds most convincing are, on this machine, among the least faithful. Fourteen methods fall in the danger zone, defined here as a proxy of at least 0.70 together with a faithfulness below 0.40. They include the entire gradient battery, the correlational neuroscience methods, linear probing, and the sparse autoencoder. The per-method extreme makes the point sharply. Vanilla saliency has the highest proxy, 0.9, yet scores faithfulness 0.0 on the position regime. Activation patching has a lower proxy, 0.5, and scores 1.0.

Here the less faithful method is the more plausible one. These figures are produced by `tools/xai_study/compare/plausibility_sensitivity.py` from the leaderboard and recorded in `compare/out/plausibility_sensitivity.csv`.


A proxy fixed by a coarse rubric raises the question of whether the contrast is an artifact of the particular numbers. The answer is that it is not. We perturb the proxy under four deterministic schemes and recompute the relationship each time. The analysis uses fixed, index-keyed offsets rather than random draws, so it re-derives bit-for-bit. The first scheme adds a fixed per-method jitter of magnitude $\sigma \in \{0.05, 0.1, 0.2\}$, alternating in sign by row. The second compresses every proxy toward the common mean by a fraction σ , which preserves the rank order of traditions. The third uses the same index-keyed offsets at $\sigma \in \{0.1, 0.2\}$, large enough to swap neighbouring traditions' ranks. The fourth drops one whole tradition at a time and recomputes on the remainder. Across all fourteen perturbations the rank correlation between faithfulness and proxy stays negative. It ranges from -0.79 to -0.40 and never approaches zero. The danger zone stays populated, with between six and fifteen methods. The per-method anchor sign holds in every scheme where both anchor methods are present, the more plausible saliency remaining the less faithful. The worst case is the leave-the-gradient-tradition-out run. It removes the most plausible and least faithful block at once, and still leaves six danger-zone methods and a rank correlation of -0.79 . Table S10 reports the full sweep. The headline contrast survives every alternative proxy assignment we tried: plausibility and faithfulness are not the same axis, and high plausibility is no guarantee of faithfulness.

Table S8 Headline-number provenance. Each number $\rightarrow$ its committed record and JSON key $\rightarrow$ the generating or aggregating script $\rightarrow$ the command that reproduces it. Values are the measured, committed numbers at this revision (consistent with REPRODUCIBILITY Section 6.5); CIs are read from `leaderboard_ci.csv`. Paths are relative to `tools/xai_study/`.

Number	Value	Record (JSON key)	Command
Faithfulness gap, position regime (causal – gradient)	0.3435	<code>faithful_demo: value; leaderboard: .. .position_regime.gap</code>	<code>faithful_demo.py</code>
Causal/interv. mean F, position regime $n=4$	0.4118 ± 0.39	<code>leaderboard: ...position_regime .causal_..._mean</code>	<code>leaderboard.py</code>
Gradient/corr. mean F, position regime $n=9$	0.0683 ± 0.07	<code>leaderboard: ...position_regime .gradient_..._mean</code>	<code>leaderboard.py</code>
Causal/interv. mean F, all regimes $n=11$	0.6811	<code>leaderboard_ci.csv: family_causal _intervention mean</code>	<code>uncertainty.py</code>
Gradient/corr. mean F, all regimes $n=14$	0.2938	<code>leaderboard_ci.csv: family_gradient _correlational mean</code>	<code>uncertainty.py</code>
Faithfulness gap, all regimes	0.3873	<code>leaderboard_ci.csv: all_regime_gap mean</code>	<code>uncertainty.py</code>
Decisive pair, position regime vs. (act. patch vs. vanilla sal.)	1.000 vs. 0.000	<code>faithful_demo: pair _contrast</code>	<code>faithful_demo.py</code>
Plausibility proxy, decisive pair	0.5 vs. 0.9	<code>faithful_demo: .. .plausibility_proxy</code>	<code>faithful_demo.py</code>
ACDC sufficiency S	0.3247	<code>leaderboard_ci.csv: ACDC S_mean</code>	<code>uncertainty.py</code>
Whole-state minimality M	0.0703	<code>leaderboard_ci.csv: A8 M_mean</code>	<code>uncertainty.py</code>
Methods scored (rows)	31	<code>leaderboard: n _methods</code>	<code>leaderboard.py</code>
Per-game records aggregated	257	<code>leaderboard: n_per_game_records _aggregated</code>	<code>leaderboard.py</code>
Naive gradient, index 0; discrete index content output	index 1.0	<code>oracle_grad_pong .json: value</code>	<code>oracle_grad.jl</code>
Exact oracle max $ \Delta \text{score} $ (Pong)	17.0	<code>oracle_pong_score .json: value</code>	<code>oracle_intervene.jl</code>

Scripts live under `tools/xai_study/compare/` (Python) and `tools/xai_study/ground_truth/` (Julia); records under the corresponding `out/` trees. JSON keys are abbreviated with “...” where nested; the full paths are in REPRODUCIBILITY Section 6.5.

A taxonomy of interpretability failure modes on the VCS, by underlying mechanism (full version — Supplement)



F = point estimate; whisker = 95% bootstrap CI over games. All numbers read from committed records; full data provenance in the Supplement. The full benchmark is 30 interpretability methods + the intervention oracle positive control = 31 rows; only the failure-exhibiting subset is drawn here.

Fig. S2 Full failure taxonomy. The full-density version of the main-text taxonomy (Fig. 6). It shows the families of interpretability failure observed on the machine, each annotated with its faithfulness against the intervention oracle (Section 3.2), reported as F with its 95% bootstrap confidence interval. The leaf values trace to `compare/out/leaderboard.json` and `leaderboard.ci.csv` via Sec. S5. Note that the gradient and correlational branches collapse on the discrete position regime, where the naive gradient is provably zero (Prop. 1).

Table S9 Plausibility-proxy rubric. The proxy assigned to each method tradition, with the documented rationale. Values are the `plausibility_proxy` field of the leaderboard rows (`compare/out/leaderboard.json`). They are tradition-level and fixed before any faithfulness score is computed, so the proxy cannot be tuned to the measured faithfulness. A higher value means an explanation that the field tends to find more convincing, not one that is more faithful. The oracle (Section 3.2) is the positive control and is not a method under test.

Tradition	Proxy	Rationale (documented basis)
gradient	0.90	Dense input-space heatmaps read as a visual account of the decision; the most widely adopted XAI outputs, routinely shown as if self-explanatory.
probing	0.85	A trained decoder that “finds” a concept reads as evidence the concept is used; high adoption, though decodability is not causal use.
correlational	0.80	Tuning curves and coupling spectra match the neuroscience idiom of a convincing single-unit account.
dim. reduction	0.75	Low-dimensional factors and dictionary atoms look like a clean summary of the representation.
intervention	0.55	Occlusion and counterfactual outputs are quantitative tables, less immediately picture-like, treated as tests rather than stories.
causal	0.50	Circuit edge lists and patching results are austere artifacts rarely mistaken for a finished narrative.
descriptive	0.30	Whole-state recording is an explicit non-explanation: it withholds nothing and reads as a dump, not an account.
oracle (control)	1.00	The Section 3.2 ground truth by construction; positive control, not a method under test.

Table S10 Proxy-sensitivity sweep. The faithfulness-versus-proxy relationship under each perturbation, from `compare/out/plausibility_sensitivity.csv`. “ ρ ” is the Spearman rank correlation between faithfulness and proxy over the methods present. “danger” is the count of high-proxy (≥ 0.70), low-faithfulness (< 0.40) methods. “anchor” is 1 when the more plausible saliency is the less faithful of the saliency/activation-patching pair, and blank where a leave-one-out run removes an anchor. The baseline rubric is the first row. Across every scheme the correlation stays negative and the danger zone stays populated.

Scheme	Param	n	ρ	Danger	Anchor
baseline	—	30	−0.64	14	1
additive jitter	$\sigma = 0.05$	30	−0.63	14	1
additive jitter	$\sigma = 0.10$	30	−0.61	14	1
additive jitter	$\sigma = 0.20$	30	−0.55	15	1
rank-preserving	$\sigma = 0.10$	30	−0.64	14	1
rank-preserving	$\sigma = 0.20$	30	−0.64	14	1
rank-jittering	$\sigma = 0.10$	30	−0.61	14	1
rank-jittering	$\sigma = 0.20$	30	−0.55	15	1
leave-out causal	—	24	−0.40	14	
leave-out correlational	—	26	−0.67	10	1
leave-out descriptive	—	29	−0.60	14	1
leave-out dim. red.	—	27	−0.64	13	1
leave-out gradient	—	20	−0.79	6	
leave-out intervention	—	25	−0.61	14	1
leave-out probing	—	29	−0.65	13	1